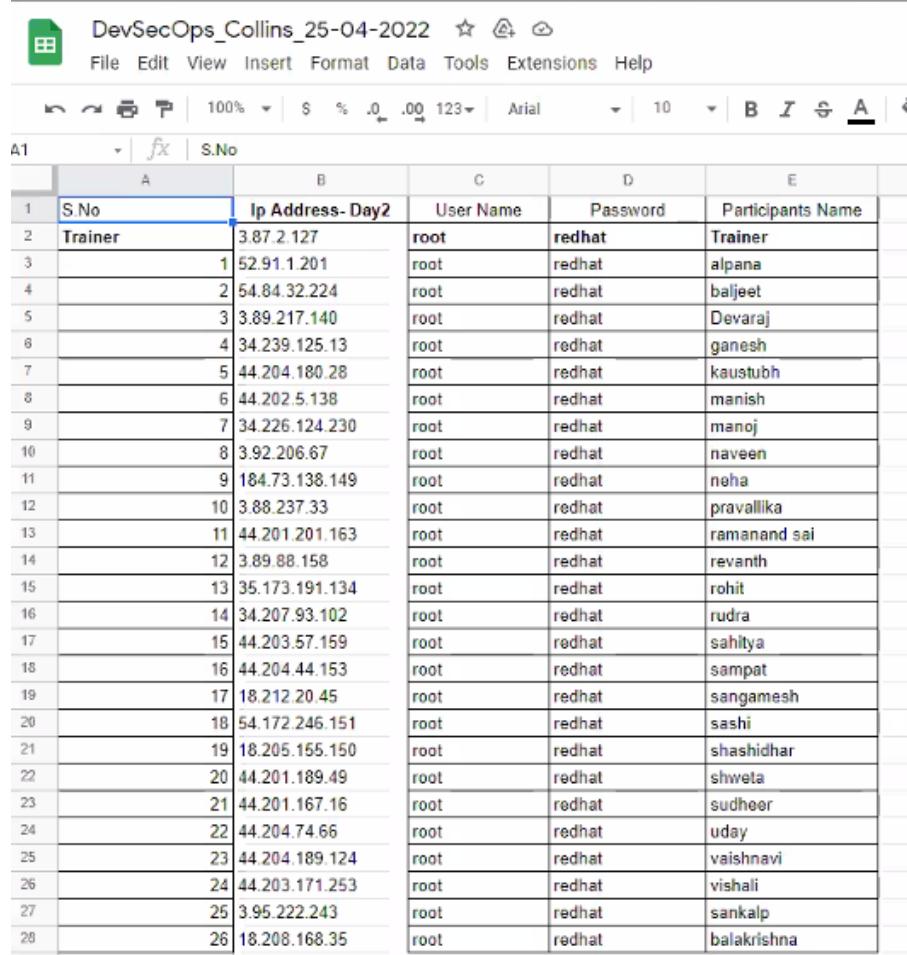


DevSecOps Training

DAY – 2

Today's IP Address details:



The screenshot shows a Microsoft Excel spreadsheet with the following data:

S.No	Ip Address- Day2	User Name	Password	Participants Name
1	3.87.2.127	root	redhat	Trainer
2	1 52.91.1.201	root	redhat	alpana
4	2 54.84.32.224	root	redhat	baljeet
5	3 3.89.217.140	root	redhat	Devaraj
6	4 34.239.125.13	root	redhat	ganesh
7	5 44.204.180.28	root	redhat	kaustubh
8	6 44.202.5.138	root	redhat	manish
9	7 34.226.124.230	root	redhat	manoj
10	8 3.92.206.67	root	redhat	naveen
11	9 184.73.138.149	root	redhat	neha
12	10 3.88.237.33	root	redhat	pravallika
13	11 44.201.201.163	root	redhat	ramanand sai
14	12 3.89.88.158	root	redhat	revanth
15	13 35.173.191.134	root	redhat	rohit
16	14 34.207.93.102	root	redhat	rudra
17	15 44.203.57.159	root	redhat	sahitya
18	16 44.204.44.153	root	redhat	sampat
19	17 18.212.20.45	root	redhat	sangamesh
20	18 54.172.246.151	root	redhat	sashi
21	19 18.205.155.150	root	redhat	shashidhar
22	20 44.201.189.49	root	redhat	shweta
23	21 44.201.167.16	root	redhat	sudheer
24	22 44.204.74.66	root	redhat	uday
25	23 44.204.189.124	root	redhat	vaishnavi
26	24 44.203.171.253	root	redhat	vishali
27	25 3.95.222.243	root	redhat	sankalp
28	26 18.208.168.35	root	redhat	balakrishna

Details of my machine:

IP Address: 44.203.171.253

Username: root

Password: redhat

Login to MobaXterm:

MobaXterm Personal Edition v22.0 •
 (SSH client, X server and network tools)

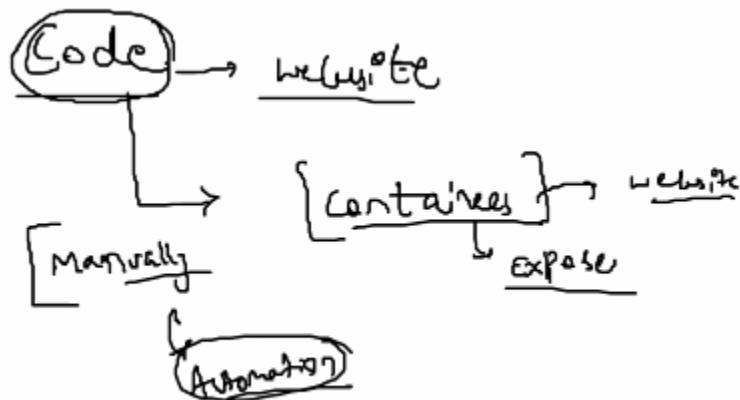
- SSH session to root@44.203.171.253
 - Direct SSH : ✓
 - SSH compression : ✓
 - SSH-browser : ✓
 - X11-forwarding : ✘ (disabled or not supported by server)

Last failed login: Mon Apr 25 13:22:36 UTC 2022 from 92.255.85.135 on ssh:notty
 There was 1 failed login attempt since the last successful login.
 Last login: Mon Apr 25 11:55:19 2022 from 183.82.31.194

Amazon Linux 2 AMI

<https://aws.amazon.com/amazon-linux-2/> [root@vishali ~]#

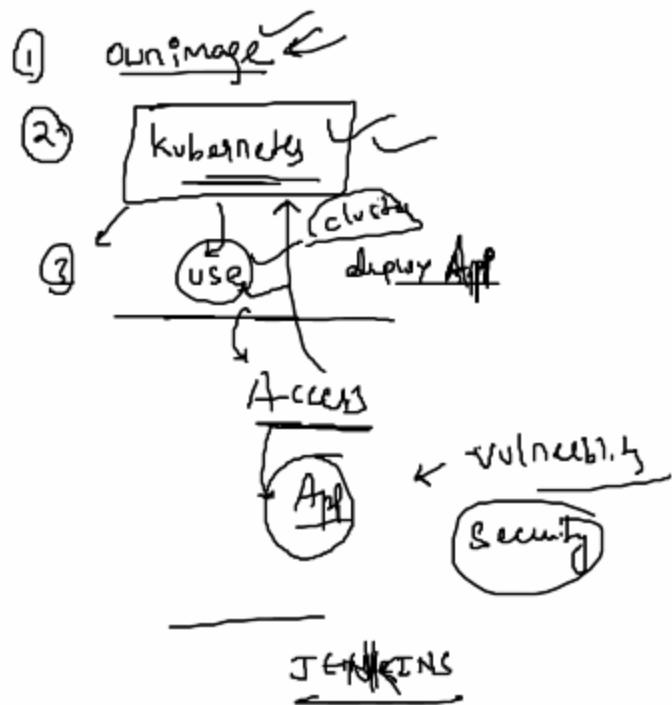
Yesterday we did the process manually



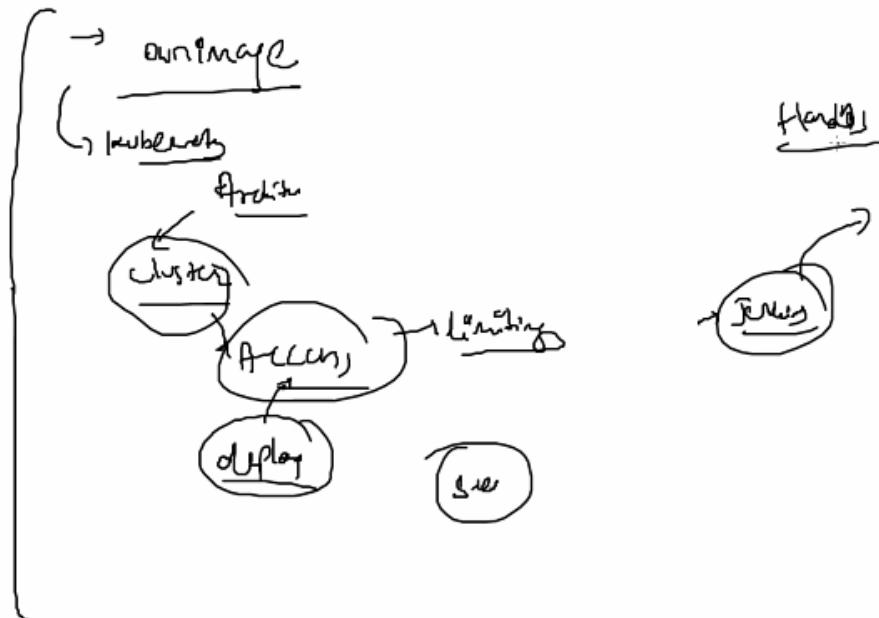
Now, we will learn to automate

Agenda for today:

- How to build the image
- Kubernetes (docker doesn't know how to manage, so we will learn this)
- How to use Kubernetes – Using this we can easily deploy our application
- How to create and access of the kubernetes cluster
- Create an application with your image – how to check the security /enhance the security/vulnerability/ what things we should take care of while deploying the application
- Jenkins



Kubernetes -> Create cluster -> Access -> Deploy -> Security

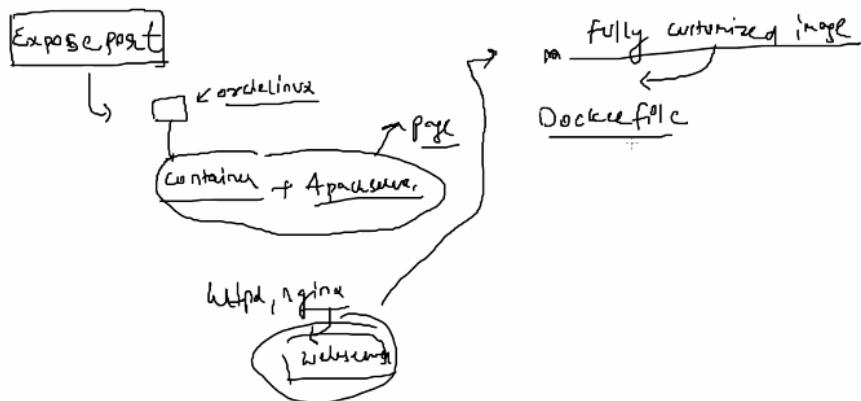


Expose port number:

We were using image oraclelinux image -> create container + Configure Apache server -> Access page.

Httpd, nginx -> These images already have the Apache server configured -> Webserver

These are fully customized image -> Dockerfile

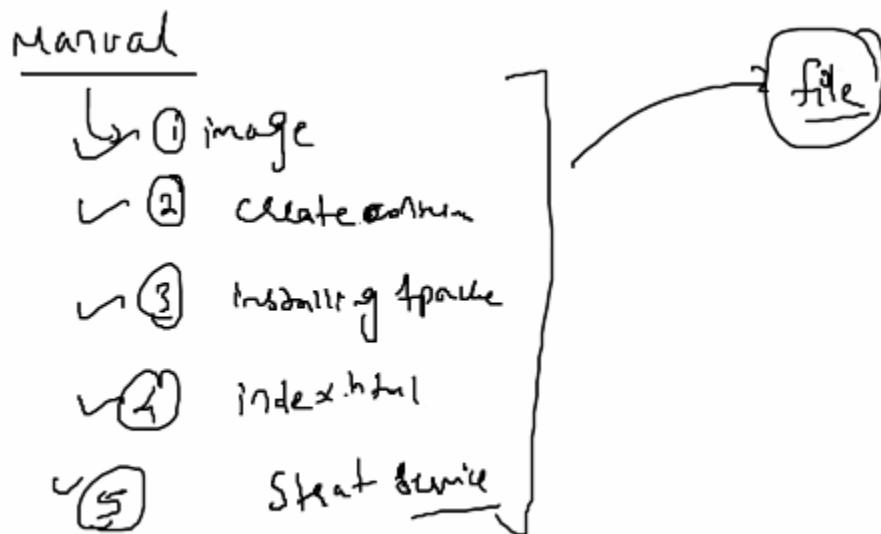


If you want to write something inside the dockerfile, what language you follow?

Programmatic language? NO

Declarative language? Yes, somewhat!

It is based on instruction.



This can be automated when we write into the Dockerfile.

Instruction:

1. FROM
2. RUN
3. RUN
4. Parent process – when the container is created by the image - Starting of the service –
CMD

Parent process can't be given by RUN.

Httpd – what is the parent process?

This is the parent process - Httpd –DFOREGROUND -> will start the service



Suggestion:

Try to create the dockerfile in the empty folder.

Check for the status of docker, if its inactive -> start the docker

```
[root@trainer ~]# systemctl status docker
● docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled; vendor preset: disabled)
  Active: inactive (dead)
    Docs: https://docs.docker.com
[root@trainer ~]# systemctl start docker
```

Create the docker file:

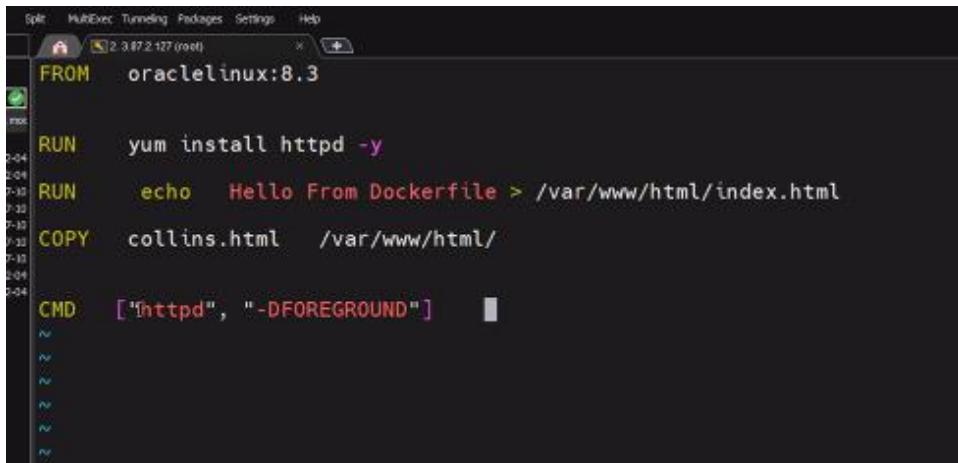
Suggestion: D - CAPS

```
[root@trainer ~]# mkdir /docker
[root@trainer ~]# cd /docker/
[root@trainer docker]# ls
[root@trainer docker]# pwd
/driver
[root@trainer docker]# vim Dockerfile
```

Dockerfile:

1. Mention the base image - FROM
2. Install the apache server - RUN
3. Create the index.html - RUN echo

4. Copy command will allow to copy local html file into the docker container location.
5. Start the service – cmd

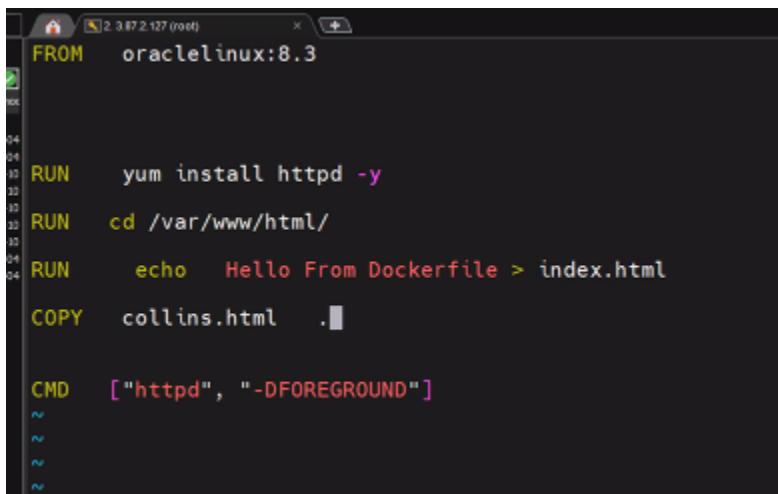


```
Split Multitab Tunneling Packages Settings Help
[2:3:872.127 (root)] x ↻
FROM oraclelinux:8.3
RUN yum install httpd -y
RUN echo Hello From Dockerfile > /var/www/html/index.html
COPY collins.html /var/www/html/
CMD ["httpd", "-DFOREGROUND"]
~
~
~
~
~
~
```

The screenshot shows a terminal window with a syntax error in the Dockerfile. The line 'COPY collins.html /var/www/html/' is highlighted in red, indicating a problem with the file path or file name.

To make it efficient:

Use cd



```
Split Multitab Tunneling Packages Settings Help
[2:3:872.127 (root)] x ↻
FROM oraclelinux:8.3
RUN yum install httpd -y
RUN cd /var/www/html/
RUN echo Hello From Dockerfile > index.html
COPY collins.html .
CMD ["httpd", "-DFOREGROUND"]
~
~
~
~
```

Use WORKDIR instead of the cd – better!

```
FROM oraclelinux:8.3

RUN yum install httpd -y
WORKDIR /var/www/html/
RUN echo Hello From Dockerfile > index.html
COPY collins.html .

CMD ["httpd", "-DFOREGROUND"]
~
~
~
~
```

Maintainer: (optional)

```
FROM oraclelinux:8.3
MAINTAINER mayank
RUN yum install httpd -y
WORKDIR /var/www/html/
RUN echo Hello From Dockerfile > index.html
COPY collins.html .

CMD ["httpd", "-DFOREGROUND"]
~
~
~
~
~
```

Expose:

```
FROM oraclelinux:8.3
MAINTAINER mayank
RUN yum install httpd -y
WORKDIR /var/www/html/
RUN echo Hello From Dockerfile > index.html
COPY collins.html .

EXPOSE 80

CMD ["httpd", "-DFOREGROUND"]
~
~
```

Save and verify:

```
[root@trainer docker]# cat Dockerfile
FROM oraclelinux:8.3

MAINTAINER mayank

RUN yum install httpd -y

WORKDIR /var/www/html/

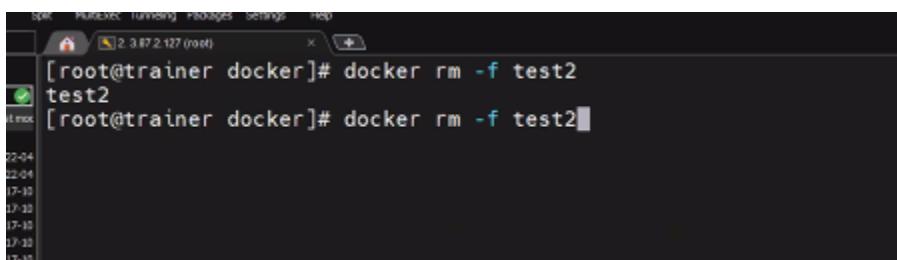
RUN echo Hello From Dockerfile > index.html

COPY collins.html .

EXPOSE 80

CMD ["httpd", "-DFOREGROUND"]
```

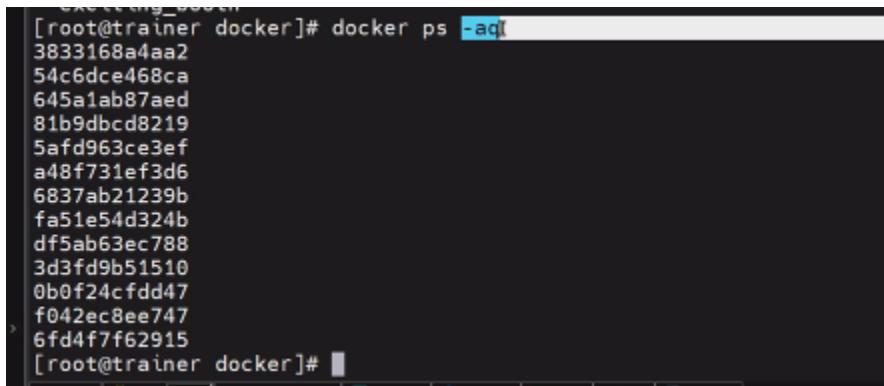
Delete the image:



The screenshot shows a terminal window with the title bar "2.3.87.2.127 (root)". The command entered is "docker rm -f test2". The output shows the container "test2" being removed successfully.

```
[root@trainer docker]# docker rm -f test2
test2
[root@trainer docker]# docker rm -f test2
```

To get all the container id:



The screenshot shows a terminal window with the title bar "executing_docker". The command entered is "docker ps -aq". The output lists numerous container IDs.

```
[root@trainer docker]# docker ps -aq
3833168a4aa2
54c6dce468ca
645a1ab87aed
81b9dbcd8219
5af963ce3ef
a48f731ef3d6
6837ab21239b
fa51e54d324b
df5ab63ec788
3d3fd9b51510
0b0f24cfdd47
f042ec8ee747
6fd4f7f62915
[root@trainer docker]#
```

All the container in one go – removed

```
[root@trainer docker]# docker rm -f $(docker ps -aq)
3833168a4aa2
54c6dce468ca
645a1ab87aed
81b9dbcd8219
5af963ce3ef
a48f731ef3d6
6837ab21239b
fa51e54d324b
df5ab63ec788
3d3fd9b51510
0b0f24cfdd47
f042ec8ee747
6fd4f7f62915
[root@trainer docker]#
```

Delete one image:

```
[root@trainer docker]# docker images
REPOSITORY          TAG      IMAGE ID      CREATED        SIZE
nginx              latest   fa5269854a5e  5 days ago    142MB
httpd              latest   c30a46771695  5 days ago    144MB
mdhack/mayanknginximage  latest   b7307762462d  8 weeks ago   142MB
hello-world        latest   feb5d9fea6a5  7 months ago   13.3kB
centos             latest   5d0da3dc9764  7 months ago   231MB
oraclelinux         8.3     816d99f0bbe8  12 months ago  224MB
mdhack/myapache    latest   ea36c5376ba4  20 months ago  166MB
mdhack/myserver    latest   e8ce7504414a  20 months ago  350MB
[root@trainer docker]# docker rmi -f hello-world
Untagged: hello-world:latest
Untagged: hello-world@sha256:10d7d58d5ebd2a652f4d93fdd86da8f265f5318c6a73cc5b6a9798ff6d2b2
Deleted: sha256:feb5d9fea6a5e9606aa995e879d862b825965ba48de054caab5ef356dc6b3412
Deleted: sha256:e07ee1baac5fae6a26f30cabfe54a36d3402f96afda318fe0a96cec4ca393359
```

All the docker images – delete at once

```
[root@trainer docker]# docker images -q
fa5269854a5e
c30a46771695
b7307762462d
5d0da3dc9764
816d99f0bbe8
ea36c5376ba4
e8ce7504414a
[root@trainer docker]# docker rmi -f $(docker images -q)
```

Build the docker image from the dockerfile:

```
[root@trainer docker]# docker build -t mayankimage .
Sending build context to Docker daemon 2.048kB
Step 1/8 : FROM oraclelinux:8.3
Step 2/8 : MAINTAINER mayank
Step 3/8 : RUN yum install httpd -y
Step 4/8 : EXPOSE 80
Step 5/8 : CMD ["httpd"]
Step 6/8 : ENTRYPOINT ["/usr/sbin/httpd"]
Step 7/8 : 
```

Package	Arch	Version	Repository	Size
httpd	x86_64	2.4.37-43.0.3.module+el8.5.0+20624+5d3b49d0.3	ol8_appstream	1.4 M

Now, the problem is we didn't create collins.html - create it

```
[root@trainer docker]# docker build -t mayankimage .
Sending build context to Docker daemon 3.072kB
Step 1/8 : FROM oraclelinux:8.3
--> 816d99f0bbe8
Step 2/8 : MAINTAINER mayank
--> Using cache
--> 47d53b1c921a
Step 3/8 : RUN yum install httpd -y
--> Using cache
--> 7b3eaa838885
Step 4/8 : WORKDIR /var/www/html/
--> Using cache
--> ae009bc9afb9
Step 5/8 : RUN echo Hello From Dockerfile > index.html
--> Using cache
--> b1a7d8609eed
Step 6/8 : COPY collins.html .
--> 9ed27d68595f
Step 7/8 : EXPOSE 80
--> Running in 9d00b77583cb
Removing intermediate container 9d00b77583cb
--> a40858ca4bf2
Step 8/8 : CMD ["httpd", "-DFOREGROUND"]
--> Running in 6dd1e6fd689a
Removing intermediate container 6dd1e6fd689a
--> 739677877265
Successfully built 739677877265
Successfully tagged mayankimage:latest
[root@trainer docker]# dock
```

Why the size has been increased? Since it is on top of oracle linux we added multiple layers

```
Successfully tagged mayankimage:latest
[root@trainer docker]# docker images
REPOSITORY TAG IMAGE ID CREATED SIZE
mayankimage latest 739677877265 8 seconds ago 393MB
oraclelinux 8.3 816d99f0bbe8 12 months ago 224MB
[root@trainer docker]#
```

Added multiple layers

```
[root@trainer docker]# docker history oraclelinux:8.3
IMAGE CREATED CREATED BY SIZE COMMENT
816d99f0bbe8 12 months ago /bin/sh -c #(nop) CMD ["/bin/bash"] 0B
<missing> 12 months ago /bin/sh -c #(nop) ADD file:8d6d5e7607cbe7af8... 224MB
[root@trainer docker]# docker history mayankimage:latest
IMAGE CREATED CREATED BY SIZE COMMENT
739677877265 46 seconds ago /bin/sh -c #(nop) CMD ["httpd" "-DFOREGROUN... 0B
a40858ca4bf2 46 seconds ago /bin/sh -c #(nop) EXPOSE 80 0B
9ed27d68595f 46 seconds ago /bin/sh -c #(nop) COPY file:f9587d7f4c8229... 21B
b1a7d8609eed About a minute ago /bin/sh -c #(nop) WORKDIR /var/www/html/ 0B
ae009bc9afb9 About a minute ago /bin/sh -c #(nop) ADD file:8d6d5e7607cbe7af8... 169MB
7b3eaa838885 About a minute ago /bin/sh -c yum install httpd -y 169MB
47d53b1c921a About a minute ago /bin/sh -c #(nop) MAINTAINER mayank 0B
816d99f0bbe8 12 months ago /bin/sh -c #(nop) CMD ["/bin/bash"] 0B
<missing> 12 months ago /bin/sh -c #(nop) ADD file:8d6d5e7607cbe7af8... 224MB
[root@trainer docker]#
```

Create the container:

```
[root@trainer docker]# docker images
REPOSITORY TAG IMAGE ID CREATED SIZE
mayankimage latest 739677877265 About a minute ago 393MB
oraclelinux 8.3 816d99f0bbe8 12 months ago 224MB
[root@trainer docker]# docker run -itd --name newcont1 -p 3333:80 mayankimage
3fc615f00ce99d4f44af8ad112b5926aa5324247db1b8365edaadf31185fd42
[root@trainer docker]#
```

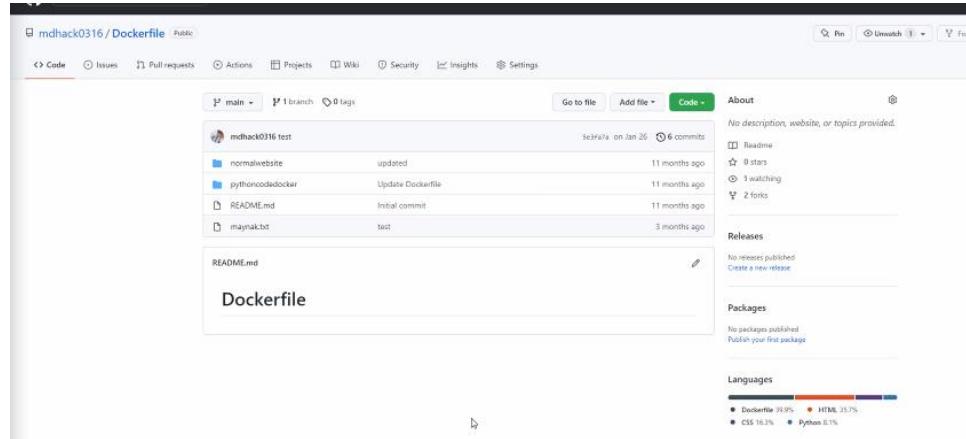
Now, you can check the website.

You will see the message written inside the index.html

If someone didn't mention the port number – random port number will be assigned!

```
[root@trainer docker]# docker run -itd -P mayankimage
1737c454adcabd5318795f1bb9db44be1b07aa44637456e090ba3eabb295da
[root@trainer docker]# docker ps -a
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS
 NAMES
1737c454adc mayankimage "httpd -DFOREGROUND" 2 seconds ago Up 1 second 0.0.0.0:49153->80/tcp, :::49153->80
0/tcp _melly_mendelev
3fcbe15f00ce mayankimage "httpd -DFOREGROUND" 2 minutes ago Up 2 minutes 0.0.0.0:3333->80/tcp, :::3333->80
tcp newcont1
[root@trainer docker]#
```

GitHub:



Here, we have a dockerfile

For downloading this repository, we need to install git

Git install - yum install git -y

git

```
[root@vishali ~]# yum install git -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core
| 3.7 kB     00:00
Resolving Dependencies
--> Running transaction check
--> Package git.x86_64 0:2.32.0-1.amzn2.0.1 will be installed
--> Processing Dependency: perl-Git = 2.32.0-1.amzn2.0.1 for package: git-2.32
-1.amzn2.0.1.x86_64
--> Processing Dependency: git-core-doc = 2.32.0-1.amzn2.0.1 for package: git-
32.0-1.amzn2.0.1.x86_64
--> Processing Dependency: git-core = 2.32.0-1.amzn2.0.1 for package: git-2.32
-1.amzn2.0.1.x86_64
--> Processing Dependency: glibc-filesystem = 27.1 for package: git-2.32.0-1.amzn2.0.1.x86_64
```

Git clone:

```
[root@trainer ~]# git clone https://github.com/mdhack0316/Dockerfile
Cloning into 'Dockerfile'...
remote: Enumerating objects: 29, done.
remote: Counting objects: 100% (29/29), done.
remote: Compressing objects: 100% (23/23), done.
remote: Total 29 (delta 3), reused 19 (delta 1), pack-reused 0
Receiving objects: 100% (29/29), 61.64 KiB | 20.55 MiB/s, done.
Resolving deltas: 100% (3/3), done.
[root@trainer ~]# ls
collins Dockerfile
[root@trainer ~]#
```

Check for the files inside it:

```
[root@trainer ~]# ls
[root@trainer ~]# cd Dockerfile/
[root@trainer Dockerfile]# s
:bash: s: command not found
[root@trainer Dockerfile]# ls
maynak.txt normalwebsite pythoncodedocker README.md
[root@trainer Dockerfile]#
```

Change the name of the docker file

```
[root@trainer normalwebsite]# mv Dockerfile mayankfile
[root@trainer normalwebsite]# docker
```

Dockerfile – not found

```
[root@trainer normalwebsite]# docker build -t example1 .
unable to prepare context: unable to evaluate symlinks in Dockerfile path: lstat /root/Dockerfile/normalwebsite: no such file or directory
[root@trainer normalwebsite]#
```

If you want to change the docker file name, then specify the -f tag:

```
[root@trainer normalwebsite]# docker build -t example1 -f mayankfile
```

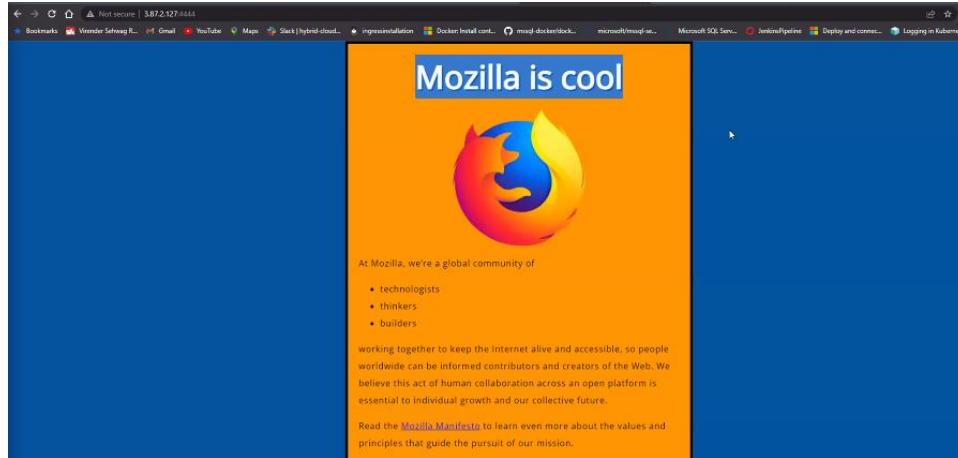
Build the image:

```
[root@trainer normalwebsite]# docker build -t example1 -f mayankfile .
Sending build context to Docker daemon 72.19kB
Step 1/5 : FROM oraclelinux:8.3
--> 816d99f0bbe8
Step 2/5 : MAINTAINER mayank123modi@gmail.com
--> Running in 038fd7fcc6fb
Removing intermediate container 038fd7fcc6fb
--> 3fd9dc24a5a5
Step 3/5 : RUN dnf install httpd -y
--> Running in 1010d1197d3f
Oracle Linux 8 BaseOS Latest (x86_64)          71 MB/s | 43 MB    00:00
```

Run the container:

```
[root@trainer normalwebsite]# docker run -itd -p 4444:80 example1  
4ea14fe899c89fe56a0d36574306d00765d6f9ddcc7c8f665643158210a3616d  
[root@trainer normalwebsite]# vim
```

Check the website:



Verify the index.html:

```
[root@trainer normalwebsite]# vim index.html  
[root@trainer normalwebsite]# docker ps
```

A screenshot of a terminal window showing the source code of the "index.html" file. The code is written in HTML and includes a DOCTYPE declaration, a head section with meta tags and links to external CSS files, and a body section containing an H1 heading, an image of the Firefox logo, a paragraph about Mozilla's global community, an ul element with three li items, and a final paragraph about the Mozilla Manifesto.

Change the content inside index.html:

Open the container in exec mode.

```
[root@trainer normalwebsite]# docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
 NAMES
4ea14fe899c8        example1          "httpd -DFOREGROUND"   42 seconds ago    Up 40 seconds     0.0.0.0:4444->80/tcp, :::4444->80/tcp
0/tcp                crazy_spence      "httpd -DFOREGROUND"   7 minutes ago     Up 6 minutes      0.0.0.0:49153->80/tcp, :::49153->80/tcp
1737c454adca        mayankimage       "httpd -DFOREGROUND"   9 minutes ago     Up 8 minutes      0.0.0.0:3333->80/tcp, :::3333->80/tcp
3fc615f0ce          mayankimage       "httpd -DFOREGROUND"   9 minutes ago     Up 8 minutes      0.0.0.0:3333->80/tcp, :::3333->80/tcp
0/tcp                newcont1
[root@trainer normalwebsite]# docker exec -it crazy_spence bash
[root@4ea14fe899c8 /]# cd /var/www/html/
[root@4ea14fe899c8 html]# vi index.html
[root@4ea14fe899c8 html]#
```

Change the content here:

```
<!DOCTYPE html>
<html>
  <head>
    <meta charset="utf-8">
    <title>My test page</title>
    <link href="http://fonts.googleapis.com/css?family=Open+Sans" rel="stylesheet">
    <link href="styles/style.css" rel="stylesheet" type="text/css">
  </head>
  <body>
    <h1>Collins is cool</h1>
    At Mozilla, we're a global community of</p>
    <ul> <!-- changed to list in the tutorial -->
      <li>technologists</li>
      <li>thinkers</li>
      <li>builders</li>
    </ul>
    <p>working together to keep the Internet alive and accessible, so people worldwide can be informed contributors and creators of the Web. We believe this act of human collaboration across an open platform is essential to individual growth and our collective future.</p>
    <p>Read the <a href="https://www.mozilla.org/en-US/about/manifesto/">Mozilla's manifesto</a> to learn about the values and principles that guide the pursuit of our mission.</p>
  </body>
</html>
```

Now you can see the content change in the website:



To see what containers are running:

```
[root@trainer normalwebsite]# docker ps -a
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
4ea14fe899c8        example1           "httpd -DFOREGROUND"   About a minute ago   Up About a minute   0.0.0.0:4444->80/tcp, ::1:4444->80/tcp
1737c454adca        mayankimage        "httpd -DFOREGROUND"   8 minutes ago      Up 7 minutes      0.0.0.0:49153->80/tcp, ::1:49153->80/tcp
3fc615f00ce         jolly_mendelev    "httpd -DFOREGROUND"   10 minutes ago     Up 9 minutes      0.0.0.0:3333->80/tcp, ::1:3333->80/tcp
[root@trainer normalwebsite]#
```

Python code:

```
[root@trainer pythoncodedocker]# cat hello.py
import time

while True:
    print("Hello all , welcome to python..!!!")
    time.sleep(3)
    print("Welcome to Mayank Containers")
    time.sleep(2)
    print("Welcome to Containers ..!!!")
    print("-----")
    time.sleep(3)

[root@trainer pythoncodedocker]# ls
Dockerfile  hello.py
[root@trainer pythoncodedocker]#
```

What is it trying to do? Change/Print the statement after few seconds

Understand the dockerfile:

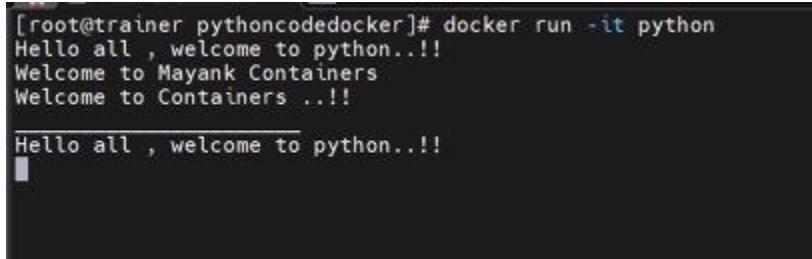
```
FROM oraclelinux:8.3
# choosing OL as docker image OS
# image we are choosing it must be present in Docker HUB
# FROM instruction will pull the image from DOCKER HUB
MAINTAINER mayank123modi@gmail.com
# image builder info (optional field)
RUN dnf install python3 -y
# DNF is an automated installer in OL
# DNF will download python and install it from Oracle registry
# RUN will be giving a temporary shell to execute INSTRUCTION
RUN mkdir /mycode
# creating directory to copy code here
COPY hello.py /mycode/hello.py
# copy data from Docker client to Docker Engine
# under that image -- that is getting build
CMD ["python3","/mycode/hello.py"]
# So this is for running code
# BUT only as parent process
# CMD is going to decide default parent process for
# this new image
```

Build the dockerfile:



```
[root@trainer pythoncodedocker]# vim Dockerfile
[root@trainer pythoncodedocker]# docker build -t python .
Sending build context to Docker daemon 3.584kB
Step 1/6 : FROM oraclelinux:8.3
--> 816d99f0bbe8
Step 2/6 : MAINTAINER mayank123modi@gmail.com
--> Using cache
--> 3fd9dc24a5a5
Step 3/6 : RUN dnf install python3 -y
--> Running in eead6edaficc
Oracle Linux 8 BaseOS Latest (x86_64)          80 MB/s | 43 MB     00:00
```

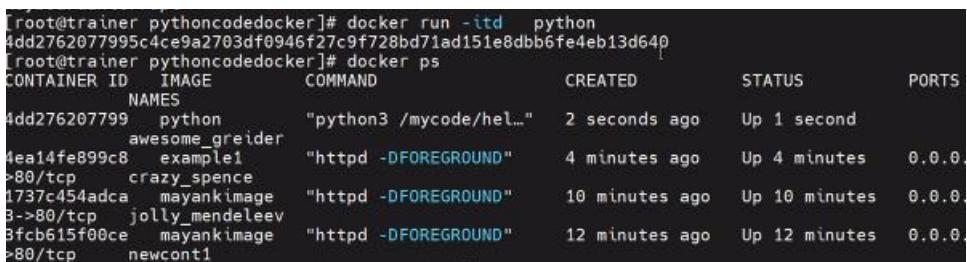
To run the container in the foreground:



```
[root@trainer pythoncodedocker]# docker run -it python
Hello all , welcome to python..!!
Welcome to Mayank Containers
Welcome to Containers ...!!

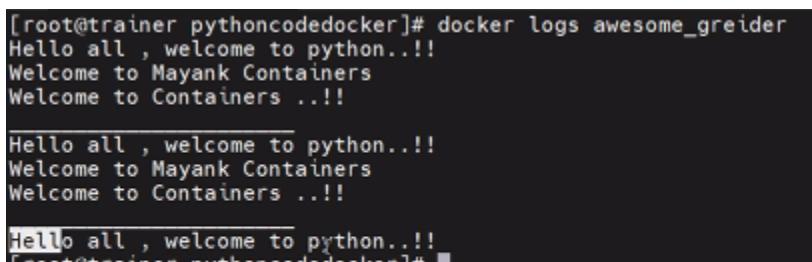
Hello all , welcome to python..!!
```

To run the container in the background:



```
[root@trainer pythoncodedocker]# docker run -itd python
4dd2762077995c4ce9a2703df0946f27c9f728bd71ad151e8dbb6fe4eb13d640
[root@trainer pythoncodedocker]# docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS
 NAMES
4dd276207799 python "python3 /mycode/hel..." 2 seconds ago Up 1 second
awesome_greider
4ea1afe899c8 example1 "httpd -DFOREGROUND" 4 minutes ago Up 4 minutes 0.0.0.
>80/tcp crazy_spence
1737c454adca mayankimage "httpd -DFOREGROUND" 10 minutes ago Up 10 minutes 0.0.0.
>80/tcp jolly_mendeleev
3fcfb615f00ce mayankimage "httpd -DFOREGROUND" 12 minutes ago Up 12 minutes 0.0.0.
>80/tcp newcont1
```

To check the logs:



```
[root@trainer pythoncodedocker]# docker logs awesome_greider
Hello all , welcome to python..!!
Welcome to Mayank Containers
Welcome to Containers ...!!

Hello all , welcome to python..!!
Welcome to Mayank Containers
Welcome to Containers ...!!

Hello all , welcome to python..!!
[root@trainer pythoncodedocker]#
```

Live log: Mention **-f** tag



```
[root@trainer pythoncodedocker]# docker logs -f awesome_greider
Hello all , welcome to python...
Welcome to Mayank Containers
Welcome to Containers ...
Hello all , welcome to python...
Welcome to Mayank Containers
Welcome to Containers ...
Hello all , welcome to python...
Welcome to Mayank Containers
Welcome to Containers ...
Hello all , welcome to python...
Welcome to Mayank Containers
Welcome to Containers ...
Hello all , welcome to python...
Welcome to Mayank Containers
Welcome to Containers ...!
```

Reference to Docker commands: (Shared by Trainer)

- 191 systemctl status docker
- 192 systemctl start docker
- 193 systemctl status docker
- 194 docker images
- 195 docker ps
- 196 docker ps -a
- 197 mkdir /docker
- 198 cd /docker/
- 199 ls
- 200 pwd
- 201 vim Dockerfile
- 202 cat Dockerfile
- 203 docker images
- 204 ls
- 205 docker ps
- 206 docker ps -a
- 207 docker rm -f test2
- 208 docker ps -a
- 209 docker rm -f fcf37b509826
- 210 docker ps -a

```
211 docker ps -aq
212 docker rm -f $(docker ps -aq)
213 docker ps -a
214 docker images
215 docker rmi -f hello-world
216 docker images
217 docker images -q
218 docker rmi -f $(docker images -q)
219 docker images
220 docker ps -a
221 cat Dockerfile
222 #docker build -t mayankimage
223 l
224 ls
225 docker build -t mayankimage .
226 ls
227 vim collins.html
228 docker build -t mayankimage .
229 docker images
230 docker history oraclelinux:8.3
231 docker history mayankimage:latest
232 docker iamges
233 docker images
234 docker run -itd --name newcont1 -p 3333:80 mayankimage
235 cat Dockerfile
236 <<X
237 vim Dockerfile
238 docker images
239 vim Dockerfile
240 docker run -itd -P mayankimage
```

```
241 docker ps -a
242 cd
243 git
244 yum install git -y
245 git clone https://github.com/mdhack0316/Dockerfile
246 ls
247 cd Dockerfile/
248 s
249 ls
250 cd normalwebsite/
251 ls
252 vim Dockerfile
253 mv Dockerfile mayankfile
254 docker build -t example1 .
255 docker build -t example1 -f mayankfile .
256 docker images
257 docker run -itd -p 4444:80 example1
258 vim index.html
259 docker ps
260 docker exec -it crazy_spence bash
261 docker ps -a
262 cd ..
263 ls
264 cd pythoncodedocker/
265 ls
266 cat hello.py
267 ls
268 vim Dockerfile
269 docker build -t python .
270 docker run -it python
```

```
271 docker run -itd python
272 docker ps
273 docker logs awesome_greider
274 docker logs -f awesome_greider
275 ls
276 cd ..
277 ls
```

TASK: 1

Q.1 create a dockerfile using as base image oraclelinux:8.3 and configure apache server and there should be 3 files index.html training.html trainer.html you need to write something inside all these files. And then create image by ur name and try to run it with any random port no. and verify whether it is accessible or not

Check for the docker service and create a folder docker:

```
[root@vishali ~]# systemctl status docker
● docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled;
  et: disabled)
    Active: inactive (dead)
      Docs: https://docs.docker.com
[root@vishali ~]# systemctl start docker
[root@vishali ~]# mkdir /docker
[root@vishali ~]# cd /docker/
[root@vishali docker]# ls
[root@vishali docker]# pwd
/docke
[root@vishali docker]# vim Dockerfile
```

Docker file:

```
[root@vishali docker]# cat Dockerfile
FROM oraclelinux:8.3
RUN yum install httpd -y
RUN echo Hello from index file > /var/www/html/index.html
RUN echo Hello from training file > /var/www/html/training.html
RUN echo Hello from trainer file > /var/www/html/trainer.html
EXPOSE 80
CMD ["httpd", "-DFOREGROUND"]
[root@vishali docker]#
```

Get MohanVTerm by subscribing to the professional edition here: <https://mohanvtterm.mohantalk.net>

Build the image:

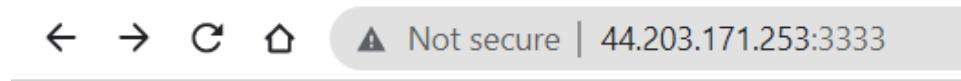
```
[root@vishali docker]# docker build -t vishaliimage .
Sending build context to Docker daemon 2.048kB
Step 1/7 : FROM oraclelinux:8.3
--> 816d99f0bbe8
Step 2/7 : RUN yum install httpd -y
--> Running in e40fe482ea23
```

Run the container:

```
[root@vishali docker]# docker run -itd --name newcont1 -p 3333:80 vishaliimage  
20e3db89d5ae82189dc59dbbab080ef9c3e5a40df38c273355152048bbcf34e  
[root@vishali docker]#
```

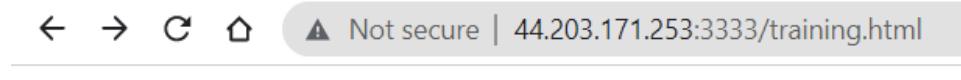
Verification:

http://44.203.171.253:3333/



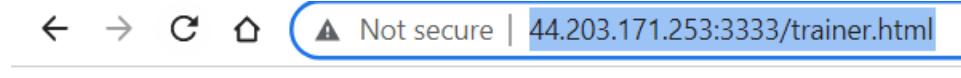
Hello from index file

http://44.203.171.253:3333/training.html



Hello from training file

<http://44.203.171.253:3333/trainer.html>



Q. 2 Use git command to clone the sae repo that I have did the URL for that is <https://github.com/mdhack0316/dockerfile> and create multiple imsing existing image and verify whether it is working fine or not

Install the git:

```
[root@vishali ~]# yum install git -y  
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd  
amzn2-core                                         | 3.7 kB     00:00  
Resolving Dependencies  
--> Running transaction check  
--> Package git.x86_64 0:2.32.0-1.amzn2.0.1 will be installed  
--> Processing Dependency: perl-Git = 2.32.0-1.amzn2.0.1 for package: git-2.32  
-1.amzn2.0.1.x86_64  
--> Processing Dependency: git-core-doc = 2.32.0-1.amzn2.0.1 for package: git-2.  
32.0-1.amzn2.0.1.x86_64  
--> Processing Dependency: git-core = 2.32.0-1.amzn2.0.1 for package: git-2.32  
-1.amzn2.0.1.x86_64  
--> Processing Dependency: libgcc_filesystem >= 0.27.1 for package: git-2.32.0-1.amzn2.0.1.x86_64
```

Clone the repo:

```
[root@vishali ~]# git clone https://github.com/mdhack0316/dockerfile
Cloning into 'dockerfile'...
remote: Enumerating objects: 29, done.
remote: Counting objects: 100% (29/29), done.
remote: Compressing objects: 100% (23/23), done.
remote: Total 29 (delta 3), reused 19 (delta 1), pack-reused 0
Receiving objects: 100% (29/29), 61.64 KiB | 20.55 MiB/s, done.
Resolving deltas: 100% (3/3), done.
```

Navigate inside normalwebsite:

```
[root@vishali ~]# ls
collins  dockerfile
[root@vishali ~]# cat dockerfile/
cat: dockerfile/: Is a directory
[root@vishali ~]# cd dockerfile/
[root@vishali dockerfile]# ls
maynak.txt  normalwebsite  pythoncodedocker  README.md
[root@vishali dockerfile]# cd normalwebsite/
[root@vishali normalwebsite]# ls
CODE_OF_CONDUCT.md  Dockerfile  images  index.html  LICENSE  README
```

Understand the docker file:

```
[root@vishali normalwebsite]# cat Dockerfile
FROM oraclelinux:8.3
# Pulling OL docker image from Docker HUB
# If not present in Docker engine
MAINTAINER mayank123modi@gmail.com
# contact me if any help required
RUN dnf install httpd -y
# DNF is an installer of OL like YUM
COPY . /var/www/html/
# C0py all the code / data from current location
# to /var/www/html/ in D0cker ENgine
CMD ["httpd","-DFOREGROUND"]
# default parent process is to start webserver
```

Build the image:

```
[root@vishali normalwebsite]# docker build -t imsing .
Sending build context to Docker daemon 72.19kB
Step 1/5 : FROM oraclelinux:8.3
--> 816d99f0bbe8
Step 2/5 : MAINTAINER mayank123modi@gmail.com
--> Running in f341ca525e83
Removing intermediate container f341ca525e83
--> 5958f6df3c8b
Step 3/5 : RUN dnf install httpd -y
--> Running in 03192839250e
Oracle Linux 8 BaseOS Latest (x86_64)          55 MB/s | 43 M
```

Create the container:

```
[root@vishali normalwebsite]# docker run -itd -p 4444:80 imsing
6553dcab51d4e8797569519baa5e3fde1b337237bafec199d4849e7d9f9db6d3
[root@vishali normalwebsite]#
```

Q.3 use docker.io/library/httpd and change the content of index file by default the content is IT WORKS, but u need to change it and try to run and verify.

Dockerfile:

```
[root@vishali task3]# cat Dockerfile
FROM docker.io/library/httpd
EXPOSE 80
CMD ["httpd", "-DFOREGROUND"]
```

Build the image:

```
[root@vishali task3]# docker build -t httpdimage .
Sending build context to Docker daemon 2.048kB
Step 1/3 : FROM docker.io/library/httpd
--> c30a46771695
Step 2/3 : EXPOSE 80
--> Running in cbcdd35dd961
Removing intermediate container cbcdd35dd961
--> 62df5835537b
Step 3/3 : CMD ["httpd", "-DFOREGROUND"]
--> Running in f78ba94ada3c
Removing intermediate container f78ba94ada3c
--> f078c9ad9969
Successfully built f078c9ad9969
Successfully tagged httpdimage:latest
```

Create the container:

```
[root@vishali task3]# docker run -itd --name httpdcont -p 3338:80 httpdimage
e96f3155a4c20c5875f3d2e9d966995db15ac01e0ba86ffb3407c6a53cb75612
```

Verify the website:



It works!

To change the content -get inside the container:

```
[root@vishali task3]# docker exec -it httpdcont bash
root@e96f3155a4c2:/usr/local/apache2# cd /var/www/html/
bash: cd: /var/www/html/: No such file or directory
root@e96f3155a4c2:/usr/local/apache2# ls
bin  build  cgi-bin  conf  error  htdocs  icons  include  logs  modules
```

Problem: Couldn't find out the index.html under the /var/www/html path

How to do we know where index.html file is present for this image?

```
[root@trainer test]# vim Dockerfile
[root@trainer test]# docker run -itd --name check1 httpd
Unable to find image 'httpd:latest' locally
latest: Pulling from library/httpd
1fe17ze4850f: Extracting [=====]
e2fa1fe9b1ec: Download complete
60dd7398e74e: Download complete
ea2ca81c6d4c: Download complete
f646c69a26ec: Download complete
```

Enter inside the container:

```
[root@trainer test]# docker exec -it check1 bash
root@63a7307c2896:/usr/local/apache2#
root@63a7307c2896:/usr/local/apache2#
root@63a7307c2896:/usr/local/apache2# pwd
/usr/local/apache2
root@63a7307c2896:/usr/local/apache2# ls
bin build cgi-bin conf error htdocs icons include logs modules
root@63a7307c2896:/usr/local/apache2#
```

Locate command

```
bin build cgi-bin conf error htdocs icons include logs modules
root@63a7307c2896:/usr/local/apache2# locate index.html
bash: locate: command not found
root@63a7307c2896:/usr/local/apache2# locate index.htm
```

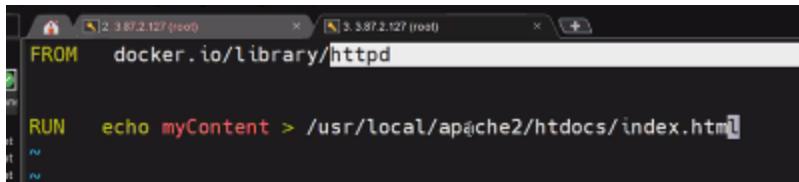
Find command:

```
root@63a7307c2896:/usr/local/apache2# c
root@63a7307c2896:/usr/local/apache2# find / index.html
```

```
root@63a7307c2896:/usr/local/apache2# find / -name index.html
find: '/proc/10/map_files': Permission denied
find: '/proc/11/map_files': Permission denied
find: '/proc/12/map_files': Permission denied
/usr/local/apache2/htdocs/index.html
root@63a7307c2896:/usr/local/apache2# ca
```

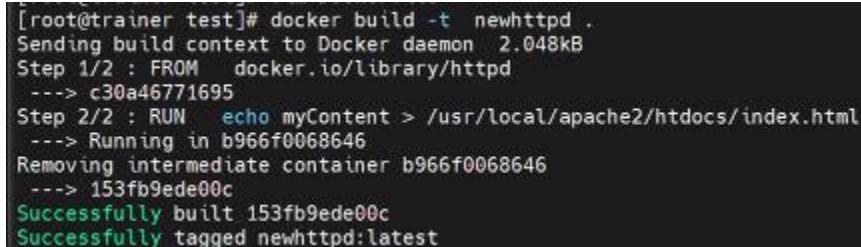
```
root@63a7307c2896:/usr/local/apache2# cat /usr/local/apache2/htdocs/index.htm
<html><body><h1>It works!</h1></body></html>
root@63a7307c2896:/usr/local/apache2#
root@63a7307c2896:/usr/local/apache2#
```

Docker file:



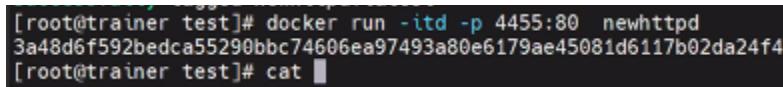
```
FROM docker.io/library/httpd
RUN echo myContent > /usr/local/apache2/htdocs/index.html
```

Build the image:



```
[root@trainer test]# docker build -t newhttpd .
Sending build context to Docker daemon 2.048kB
Step 1/2 : FROM docker.io/library/httpd
--> c30a46771695
Step 2/2 : RUN echo myContent > /usr/local/apache2/htdocs/index.html
--> Running in b966f0068646
Removing intermediate container b966f0068646
--> 153fb9ede00c
Successfully built 153fb9ede00c
Successfully tagged newhttpd:latest
```

Create the container:

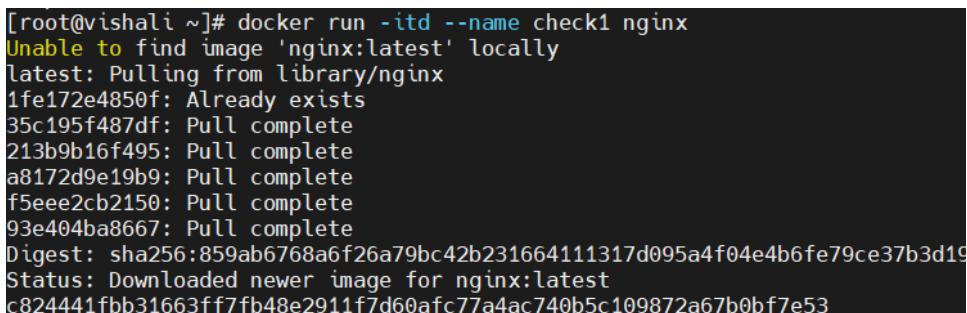


```
[root@trainer test]# docker run -itd -p 4455:80 newhttpd
3a48d6f592bedca55290bbc74606ea97493a80e6179ae45081d6117b02da24f4
[root@trainer test]# cat
```

Q.4 use nginx image and change the content of this image using docker file. and then try to expose it on any port no. and verify the same

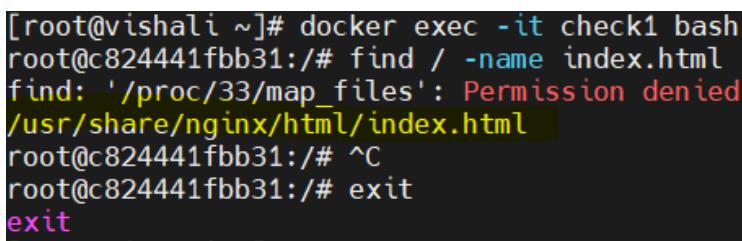
/usr/share/nginx/html/index.html

Find the location:



```
[root@vishali ~]# docker run -itd --name check1 nginx
Unable to find image 'nginx:latest' locally
latest: Pulling from library/nginx
1fe172e4850f: Already exists
35c195f487df: Pull complete
213b9b16f495: Pull complete
a8172d9e19b9: Pull complete
f5eee2cb2150: Pull complete
93e404ba8667: Pull complete
Digest: sha256:859ab6768a6f26a79bc42b231664111317d095a4f04e4b6fe79ce37b3d19
Status: Downloaded newer image for nginx:latest
c824441fbb31663ff7fb48e2911f7d60afc77a4ac740b5c109872a67b0bf7e53
```

Enter inside the container:



```
[root@vishali ~]# docker exec -it check1 bash
root@c824441fbb31:/# find / -name index.html
find: '/proc/33/map_files': Permission denied
/usr/share/nginx/html/index.html
root@c824441fbb31:/# ^C
root@c824441fbb31:/# exit
exit
```

Dockerfile:

```
[root@vishali task4]# cat Dockerfile
FROM nginx
RUN echo myContent from nginx > /usr/share/nginx/html/index.html
[root@vishali task4]#
```

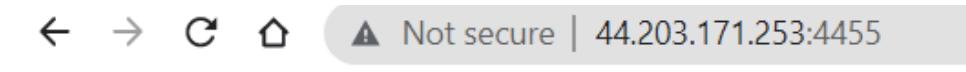
Build the image:

```
[root@vishali task4]# docker build -t newnginx .
Sending build context to Docker daemon 2.048kB
Step 1/2 : FROM nginx
--> fa5269854a5e
Step 2/2 : RUN echo myContent from nginx > /usr/share/nginx/html/index.html
--> Running in d56b7200aa9b
Removing intermediate container d56b7200aa9b
--> 20a73b7c11ab
Successfully built 20a73b7c11ab
Successfully tagged newnginx:latest
```

Run the container:

```
[root@vishali task4]# docker run -itd -p 4455:80 newnginx
2338eda4b3793cef7df7eb410d55292935acf288ce060cec8853bc35aad29de9
```

Verification:

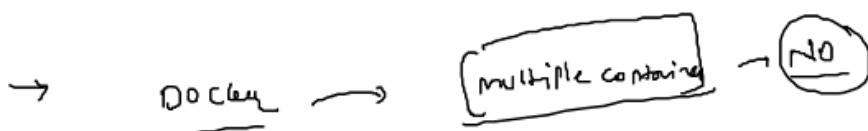


← → ⌂ ⌂ Not secure | 44.203.171.253:4455

myContent from nginx

Kubernetes:

In docker, we can't create multiple containers at once.

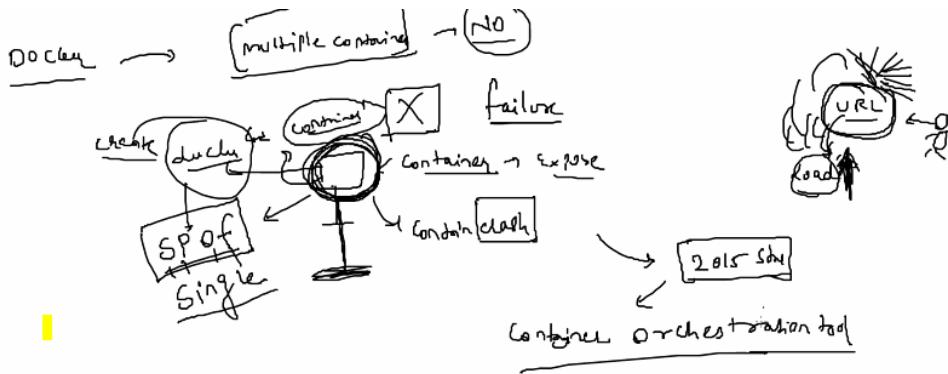


Lot of request in the URL – Load increases, time to load increases

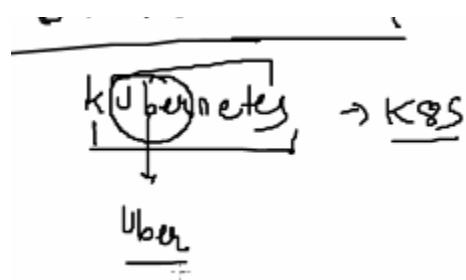
Container may crash -> Machine got crashed

Failure at the this website.

SPOF – Single point of Failure. In production, we can't create it.



Kubernetes – Orchestration Tool



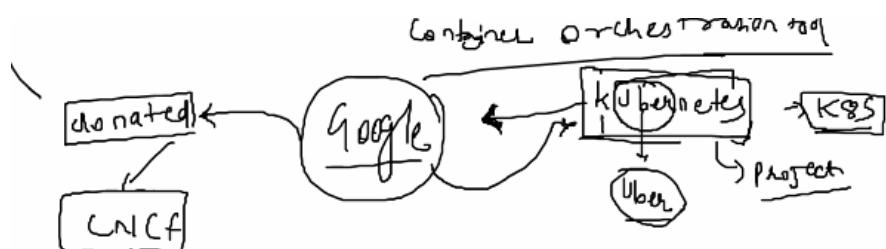
Docker is not supportive at the production.

Kubernetes

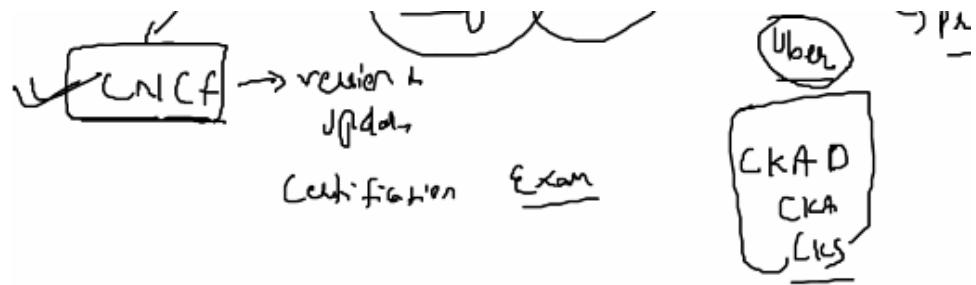
- Avoid SPOF
- Manage load
- Okay to have Multiple point of failures

Google created this project called Kubernetes in 2015 – to manage containers and remove SPOF.

Google donated to CNCF

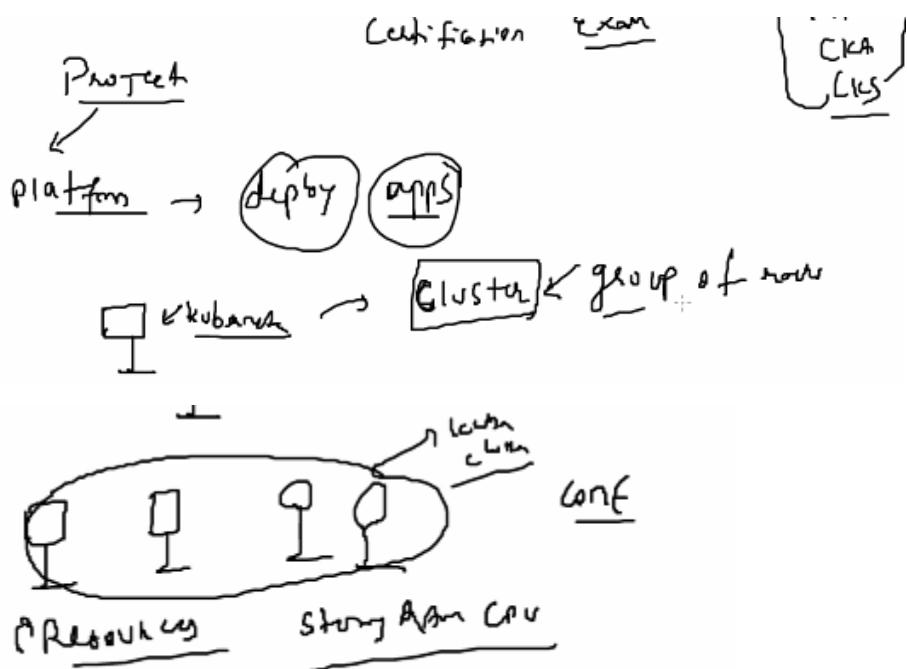


Certification exam – CKAD, CKA and CKS



Kubernetes is a tool where we can deploy our/ any time of application.

Cluster – group of machines working together.



Who will take care of the management?

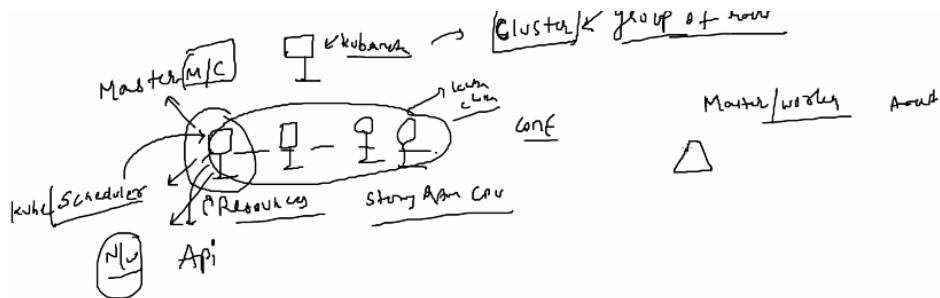
In which system the process run?

All these are taken care of Kubernetes

We will tell this machine is a master.

Who will schedule the process? kube scheduler

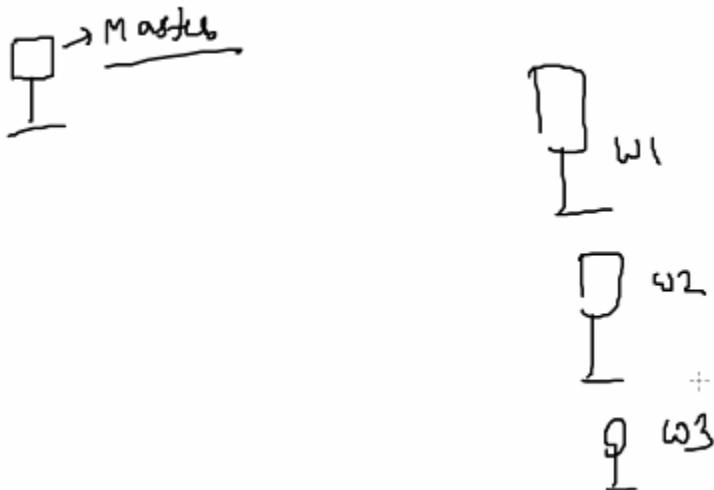
Master/Worker architecture:



Example:

New 4 machines is taken

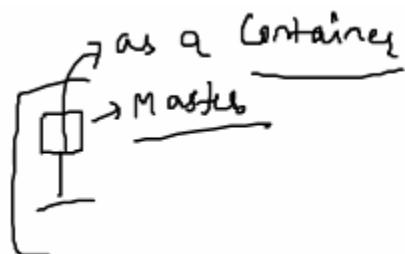
One machine will work as a master machine and remaining machines will be the worker.



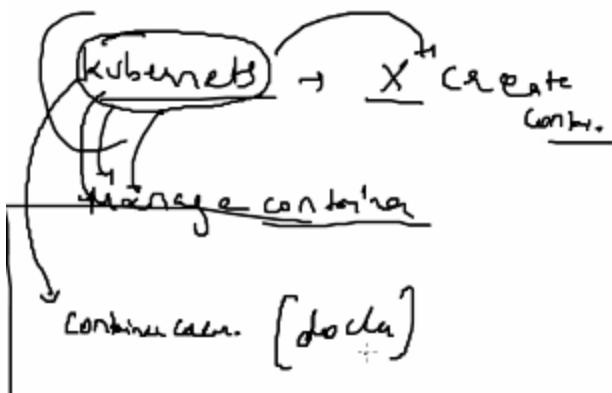
After deploying the application. One is called as master amchine and others are worker machines.

How will the architecture work?

Everything will run on the master machine as a container.

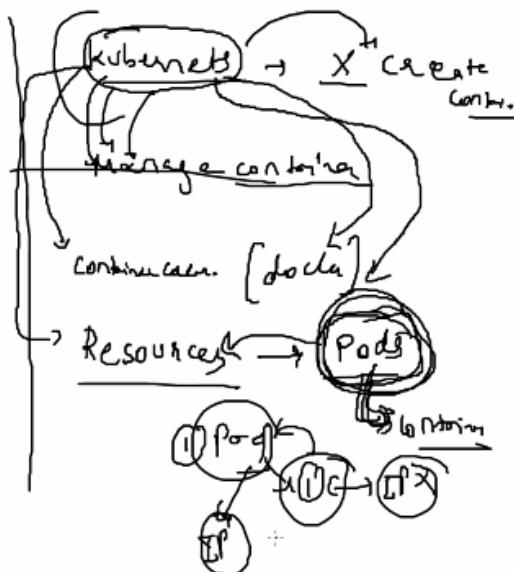


Kubernetes – won't help us to create the container. Only help us in managing the container.



Only docker will create the container. Kubernetes will have multiple resources inside it. The smallest resources in the kubernetes are called Pods. It just the virtual entity not a physical one.

Pod – resource help to create a container inside the Pod. It will get the IP address(own)



If I want to create a POD? How will I do it?

Use some commands and request will go to the master components.

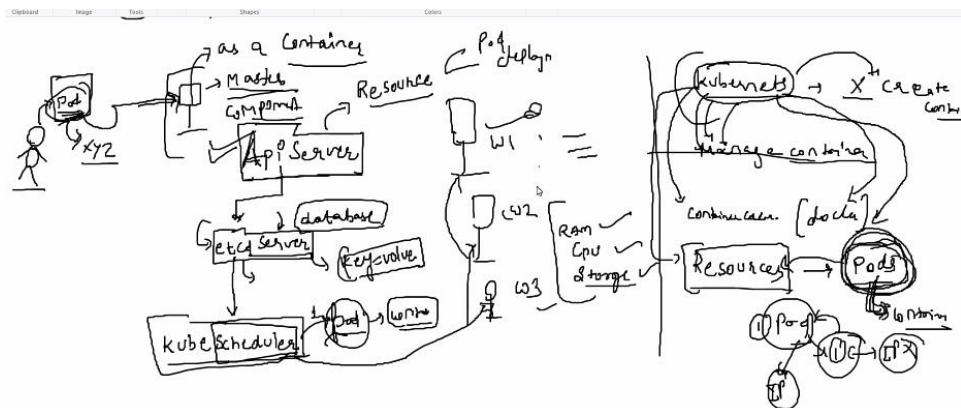
Which components? API server. API server is the one who is providing the resources like deployments and many more.

It will send my resources to new components called – etcd server

Why is my request going to etcd server? Because etcd server is the database of the kubernetes cluster. It has each and single information about the server in the form of key and value pair.

Why is my POD request going to etcd server? It will verify the pod name – already exists/running or not.

They will again send request to kube scheduler: This will assign task to each machine.



If there is a Pacific Ocean – multiple ships are running? Who will maintain the details of ships?
Captain

Like similar, for every machine we will have captain called kubelet.

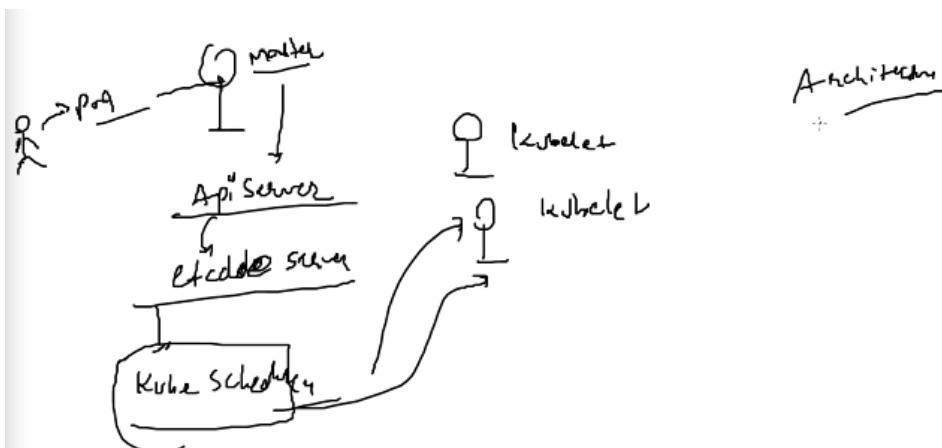
Schedule the tasks randomly in the containers.

If the worker 2 goes down, then what will happen? All the pods will be evicted from the worker 2. Pod will create another machine and manages it.

Flow:

User -> Request -> Pod-> Master> API server -> etcd server -> kubescheduler -> Kubelet

Architecture:



On what basis, will the scheduling happen? RAM, storage. RAM will take the top priority.



How to create a Kubernetes cluster?

Why am I using AWS? We need multiple machines to create cluster – so using cloud platform.

Creating the machines:

First Create EC2 instance:

- 16GB RAM in the Master
- Configure instance details
- Add storage – 40 GB
- Add tags – no need
- Configure Security group
- Review instance Launch

Change the name to use it later. Node1, node2, node 3.....and so on

Instances (1/6) Info										
	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
<input type="checkbox"/>	master	i-09727994330194ada	Running	t2.large	-	No alarms	us-east-2a	ec2-13-59-148-91.us-e...	13.59.148.91	-
<input type="checkbox"/>	node1	i-067df214d146d7179	Running	t2.large	-	No alarms	us-east-2a	ec2-3-128-78-38.us-e...	3.128.78.38	-
<input type="checkbox"/>	node2	i-08694f840dbe4519d	Running	t2.large	-	No alarms	us-east-2a	ec2-3-135-197-110.us...	3.135.197.110	-
<input type="checkbox"/>	node3	i-087eeb6537c16698	Running	t2.large	-	No alarms	us-east-2a	ec2-3-144-46-50.us-e...	3.144.46.50	-
<input type="checkbox"/>	node4	i-0463dccb78a8e6da	Running	t2.large	-	No alarms	us-east-2a	ec2-18-217-22-20.us-e...	18.217.22.20	-
<input checked="" type="checkbox"/>	node5	i-058699af04429d1c3	Running	t2.large	-	No alarms	us-east-2a	ec2-3-144-45-214.us-e...	3.144.45.214	-

We need to give elastic IP:

What is an Elastic IP?

An Elastic IP address is a **reserved public IP address that you can assign to any EC2 instance in a particular region, until you choose to release it.**

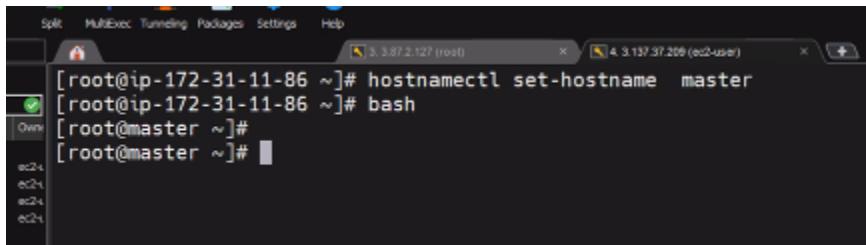
Static IP – Even if you reboot it. This IP will remain same.

Instances (1/6) Info										
	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
<input type="checkbox"/>	master	i-09727994330194ada	Running	t2.large	Initializing	No alarms	us-east-2a	ec2-3-137-209.us-e...	3.137.209	137.209
<input type="checkbox"/>	node1	i-067df214d146d7179	Running	t2.large	Initializing	No alarms	us-east-2a	ec2-3-128-78-38.us-e...	3.128.78.38	-
<input type="checkbox"/>	node2	i-08694f840dbe4519d	Running	t2.large	Initializing	No alarms	us-east-2a	ec2-3-135-197-110.us...	3.135.197.110	-
<input type="checkbox"/>	node3	i-087eeb6537c16698	Running	t2.large	Initializing	No alarms	us-east-2a	ec2-3-144-46-50.us-e...	3.144.46.50	-
<input type="checkbox"/>	node4	i-0463dccb78a8e6da	Running	t2.large	Initializing	No alarms	us-east-2a	ec2-18-217-22-20.us-e...	18.217.22.20	-
<input checked="" type="checkbox"/>	node5	i-058699af04429d1c3	Running	t2.large	Initializing	No alarms	us-east-2a	ec2-3-144-45-214.us-e...	3.144.45.214	-

Access each machine:

Connect to one machine and Assign as the master

Master:

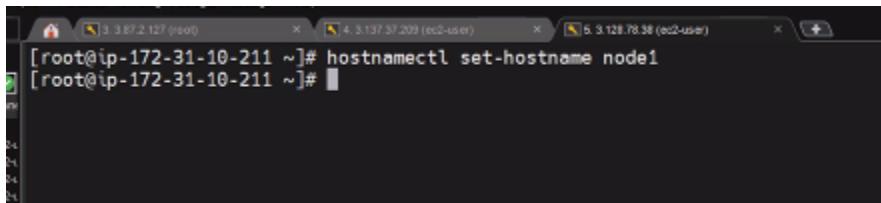


```
[root@ip-172-31-11-86 ~]# hostnamectl set-hostname master
[root@ip-172-31-11-86 ~]# bash
[root@master ~]#
[root@master ~]#
```

Set the hostname for reference purposes.

Connect to another machine:

Node1:

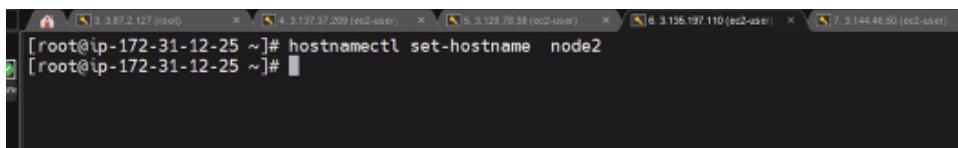


```
[root@ip-172-31-10-211 ~]# hostnamectl set-hostname node1
[root@ip-172-31-10-211 ~]#
```

Set the hostname for reference purposes.

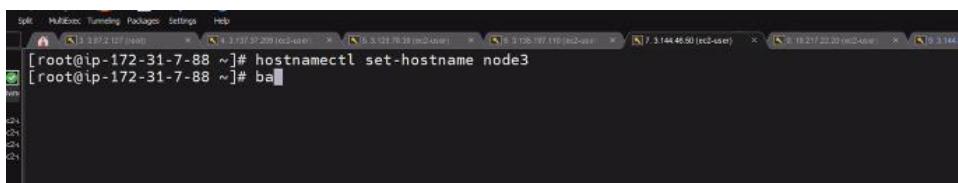
Access each machine one by one and give the hostname

Node 2:



```
[root@ip-172-31-12-25 ~]# hostnamectl set-hostname node2
[root@ip-172-31-12-25 ~]#
```

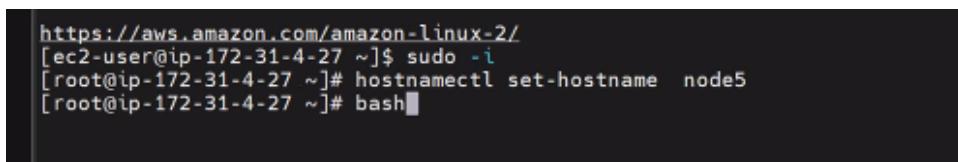
Node 3:



```
[root@ip-172-31-7-88 ~]# hostnamectl set-hostname node3
[root@ip-172-31-7-88 ~]# ba
```

Node 4:

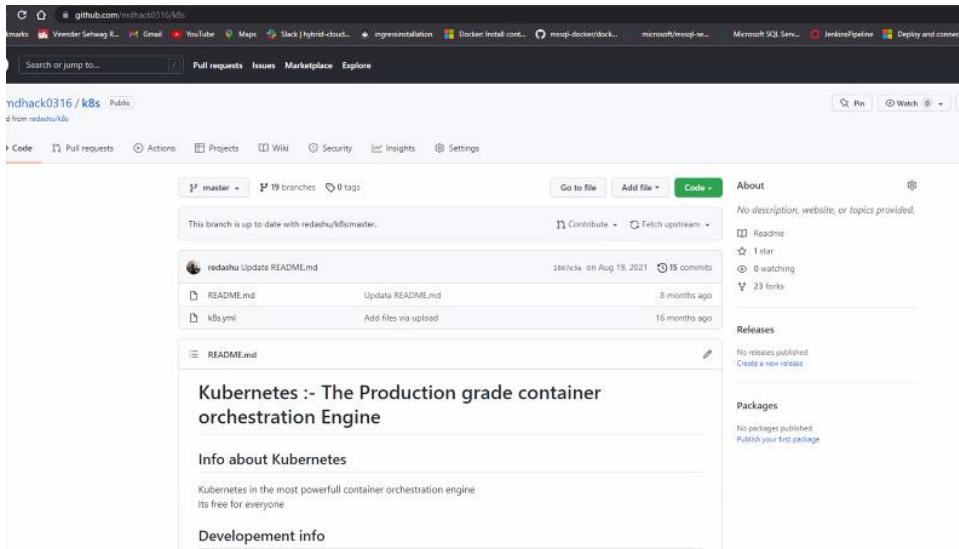
Node 5:



```
https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-172-31-4-27 ~]$ sudo -i
[root@ip-172-31-4-27 ~]# hostnamectl set-hostname node5
[root@ip-172-31-4-27 ~]# bash
```

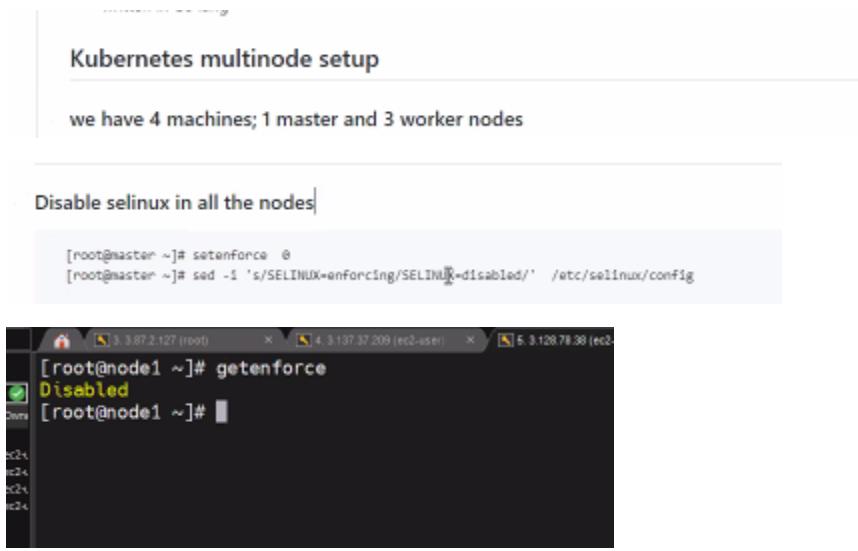
We have total 6 machines. 1 – master machine and other are worker machine.

Go to this link:



The screenshot shows a GitHub repository page for 'k8s' by 'redashu'. The repository has 19 branches and 0 tags. The README.md file contains the text 'Kubernetes :- The Production grade container orchestration Engine'. The repository has 1 star, 0 watching, and 23 forks. There are sections for 'About', 'Releases', and 'Packages'.

Create a cluster:



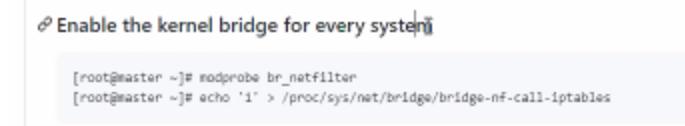
The terminal session shows the configuration of a Kubernetes multinode setup. It includes commands to disable SELinux and check node status.

```
[root@master ~]# setenforce 0
[root@master ~]# sed -i 's/SELINUX=enforcing/SELINUXTYPE=disabled/' /etc/selinux/config
```

```
[root@node1 ~]# getenforce
Disabled
[root@node1 ~]#
```

Check in every machine whether it is disabled. If not use the above command.

Step 2:



The terminal session shows the enablement of the kernel bridge for every system.

```
[root@master ~]# modprobe br_netfilter
[root@master ~]# echo '1' > /proc/sys/net/bridge/bridge-nf-call-iptables
```

Why do we need this bridge? WE are going to build a kubernetes cluster with 6 machine. Kubernetes system will create a container. We have to take care of networking part. This kernel module will take care of it.

```
Disabled
[root@node1 ~]# modprobe br_netfilter
```

This will create network in every machine. Run it in every machine.

Step 3:

[Disable the swap](#)

```
[root@master ~]# swapoff -a
```

What is swap memory? In your machine you can have extra RAM. If RAM which are not using given to swap memory. This is not supported in Kubernetes. So, we are disabling it.

```
[root@master ~]# swapoff -a
[root@master ~]#
```

In every machine run this command.

Step 4:

What software we need?

```
[root@master ~]# swapoff -a
[root@master ~]# <<X
>
>     docker -- yes
>     kubeadm - yes
>     kubelet -- yes
>
> kubectl
```

All these software are required to create the kubernetes cluster.

We can't install everything directly in linux:

```
[root@master ~]# yum install kubeadm kubelet docker kubectl -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core
No package kubeadm available.
No package kubelet available.
No package kubectl available.
Resolving Dependencies
--> Running transaction check
--> Package docker.x86_64 0:20.10.7-5.amzn2 will be installed
--> Processing Dependency: runc >= 1.0.0 for package: docker-20.10.7-5.amzn2.x86_64
--> Processing Dependency: libcgroup >= 0.40.rc1-5.15 for package: docker-20.10.7-5.amzn2.x86_64
--> Processing Dependency: containerd >= 1.3.2 for package: docker-20.10.7-5.amzn2.x86_64
--> Processing Dependency: pigz for package: docker-20.10.7-5.amzn2.x86_64
--> Running transaction check
--> Package containerd.x86_64 0:1.4.6-8.amzn2 will be installed
--> Package libcgroup.x86_64 0:0.41-21.amzn2 will be installed
--> Package pigz.x86_64 0:2.3.4-1.amzn2.0.1 will be installed
--> Package runc.x86_64 0:1.0.0-2.amzn2 will be installed
--> Finished Dependency Resolution

Dependencies Resolved

=====
| Package           | Arch      | Version        | Repository |
=====
Installing:
```

Copy this in every machine

```
root@node2 ~]# swapoff -a
root@node2 ~]# cat <<EOF >/etc/yum.repos.d/kube.repo
[kube]
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
gpgcheck=0
EOF
root@node2 ~]# yum install kubeadm kubelet docker kubectl -y
```

Then run this command:

```
[root@master ~]# yum install kubeadm kubelet docker kubectl -y
[...]
[root@master ~]# systemctl status docker
● docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled; vendor preset: disabled)
  Active: inactive (dead)
    Docs: https://docs.docker.com
[root@master ~]# systemctl start docker
```

Docker info

```
Storage Driver: overlay2
 Backing Filesystem: xfs
 Supports d_type: true
 Native Overlay Diff: true
 userxattr: false
 Logging Driver: json-file
 Cgroup Driver: cgroupfs
 Cgroup Version: 1
 Plugins:
  Volume: local
  Network: bridge host ipvlan macvlan null overlay
  Log: awslogs fluentd gcplogs gelf journald json-file local logentries splunk
 Swarm: inactive
 Runtimes: io.containerd.runc.v2 io.containerd.runtime.v1.linux runc
 Default Runtime: runc
 Init Binary: docker-init
 containerd version: d71fcfd7d8303cbf684402823e425e9dd2e99285d
 runc version: 84113eeef6fc27af1b01b3181f31bbaf708715301
 init version: de40ad0
 Security Options:
  seccomp
    Profile: default
```

Kubernetes doesn't support cgroupfs driver. So, change it to system driver

```
[root@master ~]# cat <<X >/etc/docker/daemon.json
> {
>   "exec-opts": ["native.cgroupdriver=systemd"]
> }
> X
[root@master ~]# cat /etc/docker/daemon.json
```

```
[root@master ~]# cat /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}

[root@master ~]#
```

```
[root@master ~]# cat /etc/docker/daemon.json
[root@master ~]# cat /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}

[root@master ~]#
```

Reload the daemon and restart the docker:

```
[root@master ~]# systemctl daemon-reload
[root@master ~]# systemctl restart docker
```

Docker info command:

```
Running: 0
Paused: 0
Stopped: 0
Images: 0
Server Version: 20.10.7
Storage Driver: overlay2
Backing Filesystem: xfs
Supports d_type: true
Native Overlay Diff: true
userxattr: false
Logging Driver: json-file
Cgroup Driver: systemd
Cgroup Version: 1
Plugins:
  Volume: local
  Network: bridge host ipvlan macvlan null
  Log: awslogs fluentd gcplogs gelf journal
Swarm: inactive
Runtimes: io.containerd.runtime.v1.linux
Default Runtime: runc
Init Binary: docker-init
containerd version: d71fc7d8303cbf6844028
runc version: 84113eef6fc27af1b01b3181f31b
```



```
[root@node1 ~]# systemctl daemon-reload ; systemctl restart docker ; systemctl enable kubeadm
```

This command you have to run in every single machine.

Above configuration needs to be completed in every machine.

Now we have to decide which machine is master and worker?

```
[root@master ~]# kubeadm init
```

It will be initialized on private IP address – we won't access it

So, mention the Public IP address

```
[root@master ~]# kubeadm init --apiserver-cert-extra-sans=3.137.37.209 --apiserver-advertise-address=0.0.0.0
```

When you create a kube cluster - > It will create a pod in the background – Mention the any private IP range for pod.

```
[root@master ~]# kubeadm init --apiserver-cert-extra-sans=3.137.37.209 --apiserver-advertise-address=0.0.0.0  
--pod-network-cidr=192.168.0.0/24
```

This command will only run on the master machine.

It takes 2 – 3 minutes – depends on the internet speed. Since it downloads multiple images.

```
[root@master ~]# kubeadm init --apiserver-cert-extra-sans=3.137.37.209 --apiserver-advertise-address=0.0.0.0  
--pod-network-cidr=192.168.0.0/24  
[init] Using Kubernetes version: v1.23.6  
[preflight] Running pre-flight checks  
[WARNING ServiceDocker]: docker service is not enabled, please run 'systemctl enable docker.service'  
[WARNING FileExisting-tc]: tc not found in system path  
[WARNING Hostname]: hostname "master" could not be reached  
[WARNING Hostname]: hostname "master": lookup master on 172.31.0.2:53: no such host  
[preflight] Pulling images required for setting up a Kubernetes cluster  
[preflight] This might take a minute or two, depending on the speed of your internet connection  
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
```

Public IP Address:

```
[preflight] Pulling images required for setting up a Kubernetes cluster  
[preflight] This might take a minute or two, depending on the speed of your internet connection  
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'  
[certs] Using certificate-dir folder "/etc/kubernetes/pki"  
[certs] Generating "ca" certificate and key  
[certs] Generating "apiserver" certificate and key  
[certs] apiserver serving cert is signed for DNS names [kubernetes kubernetes.default kubernetes.default.svc.  
es.default.svc.cluster.local master] and IPs [10.96.0.1 172.31.11.86 3.137.37.209]  
[certs] Generating "apiserver-kubelet-client" certificate and key  
[certs] Generating "front-proxy-ca" certificate and key  
[certs] Generating "front-proxy-client" certificate and key  
[certs] Generating "etcd/ca" certificate and key  
[certs] Generating "etcd/server" certificate and key  
[certs] etcd/server serving cert is signed for DNS names [localhost master] and IPs [172.31.11.86 127.0.0.1]  
[certs] Generating "etcd/peer" certificate and key  
[certs] etcd/peer serving cert is signed for DNS names [localhost master] and IPs [172.31.11.86 127.0.0.1 ...]
```

Created successfully:

```
[mark-control-plane] Marking the node master as control-plane by adding the taints [node-role.kubernetes.io  
Schedule]  
[bootstrap-token] Using token: 9hir34.msvosn8vdfkh01p3  
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles  
[bootstrap-token] configured RBAC rules to allow Node Bootstrap tokens to get nodes  
[bootstrap-token] configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to  
rm certificate credentials  
[bootstrap-token] configured RBAC rules to allow the csapprover controller automatically approve CSRs fro  
otstrap Token  
[bootstrap-token] configured RBAC rules to allow certificate rotation for all node client certificates in  
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace  
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certific  
[addons] Applied essential addon: CoreDNS  
[addons] Applied essential addon: kube-proxy  
  
Your Kubernetes control-plane has initialized successfully!  
  
To start using your cluster, you need to run the following as a regular user:  
  
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config  
  
Alternatively, if you are the root user, you can run:  
  
export KUBECONFIG=/etc/kubernetes/admin.conf  
  
You should now deploy a pod network to the cluster.  
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:  
https://kubernetes.io/docs/concepts/cluster-administration/addons/  
  
Then you can join any number of worker nodes by running the following on each as root:  
  
kubeadm join 172.31.11.86:6443 --token 9hir34.msvosn8vdfkh01p3 \  
--discovery-token-ca-cert-hash sha256:d7cd2db64a56153d5b0c5a0cb62b530f7190e081012ab20594672d30604
```

Authentication file

```

  Loading kubeconfig file /etc/kubernetes/admin.conf
Your Kubernetes control-plane has initialized successfully!
To start using your cluster, you need to run the following as a regular user:
  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config
Alternatively, if you are the root user, you can run:
  export KUBECONFIG=/etc/kubernetes/admin.conf
You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/
Then you can join any number of worker nodes by running the following on each as root:
  kubeadm join 172.31.11.86:6443 --token 9hir34.msvosn8vdfkh0ip3 \
    --discovery-token-ca-cert-hash sha256:d7cd2db64e56153d5be65a0cbd62b530f7190e081012ab20594672d3060f26c0
[root@master ~]#

```

We need to move this file as the kubectl should know. It will only read from specific path.

```

[root@master ~]# kubectl get nodes
The connection to the server localhost:8080 was refused - did you specify the right host or port?
[root@master ~]# 
[root@master ~]# kubectl get nodes
The connection to the server localhost:8080 was refused - did you specify the right host or port?
[root@master ~]# pwd
/root
[root@master ~]# mkdir .kube
[root@master ~]# cp /etc/kubernetes/admin.conf .kube/config
[root@master ~]# kubectl get nodes
NAME      STATUS   ROLES            AGE     VERSION
master    NotReady control-plane,master   97s    v1.23.6
[root@master ~]#

```

No worker machines are added here. To add the worker machine, copy this command in every single machine.

```

https://172.31.11.86:6443/api/v1/nodes?labelManager=kubernetes&labelSelector=&resourceVersion=0
Then you can join any number of worker nodes by running the following on each as root:
  kubeadm join 172.31.11.86:6443 --token 9hir34.msvosn8vdfkh0ip3 \
    --discovery-token-ca-cert-hash sha256:d7cd2db64e56153d5be65a0cbd62b530f7190e081012ab20594672d3060f26c0
[root@master ~]# cat /etc/kubernetes/admin.conf
[root@master ~]# cat /etc/kubernetes/admin.conf
apiversion: vi

```

Node 1, node2,

```

[root@node1 ~]# kubeadm join 172.31.11.86:6443 --token 9hir34.msvosn8vdfkh0ip3 \
>   --discovery-token-ca-cert-hash sha256:d7cd2db64e56153d5be65a0cbd62b530f7190e081012ab20594672d3060f26c0
[preflight] Running pre-flight checks
[WARNING Service-Docker]: docker service is not enabled, please run 'systemctl enable docker.service'
[WARNING FileExisting-tc]: tc not found in system path
[WARNING Hostname]: hostname "node1" could not be reached
[WARNING Hostname]: hostname "node1": lookup node1 on 172.31.0.2:53: no such host

```

Now cluster is ready!

```

[root@master ~]# kubectl get nodes
NAME      STATUS   ROLES            AGE     VERSION
master    NotReady control-plane,master   4m17s   v1.23.6
node1    NotReady <none>           29s    v1.23.6
node2    NotReady <none>           33s    v1.23.6
node3    NotReady <none>           24s    v1.23.6
node4    NotReady <none>           21s    v1.23.6
node5    NotReady <none>           5s     v1.23.6
[root@master ~]#

```

Machine status is not yet ready! Network configuring is not ready – best tool is calico networking – automating make them ready

Ad - https://try.digitalocean.com/kubernetes +

Run Kubernetes on DigitalOcean - Designed for Developers

Managed **Kubernetes** Designed for You and Your Small Business. Get Started w/ a Free Credit. Provision and deploy to your **Kubernetes** cluster in minutes. Flexible API. High-CPU Servers. Highly Scalable. Premium Intel w/ NVMe SSD. Free Monitoring & Alerts.

Kubernetes CI/CD Pipeline

Easy App Deployment & Updates. Accelerate Your Release Life-Cycle.

Tutorial Library

Integration or Development Issues? Follow One of Our 3500+ Tutorials.

Ad - https://www.sumologic.com +

K8s Monitoring - Get Started For Free - SumoLogic.com

Get Instant Insight Into The Health and Performance of **Kubernetes**. Get Started For Free. Faster Troubleshooting with Integrated Logs, Metrics and Traces. Start Your Free Trial. Native AWS Integrations. Big Data for Real-Time IT. Unified Logs & Metrics.

Install Calico with Kubernetes API datastore, 50 nodes or less

1. If you are using pod CIDR 192.168.0.0/16, skip to the next step. ...
2. Customize the manifest as necessary.
3. Apply the manifest using the following command: \$ kubectl apply -f calico.yaml

Copy this command:

many clients so using Typha is redundant and not recommended

- Install Calico with Kubernetes API datastore, 50 nodes or less
- Install Calico with Kubernetes API datastore, more than 50 nodes
- Install Calico with etcd datastore

Install Calico with Kubernetes API datastore, 50 nodes or less

1. Download the Calico networking manifest for the Kubernetes API datastore

```
$ curl https://projectcalico.docs.tigera.io/manifests/calico.yaml -O
```

Copy

2. If you are using pod CIDR 192.168.0.0/16, skip to the next step. If you are using a different pod CIDR with kubeadm, no changes are required - Calico will automatically detect the CIDR based on the running configuration. For other platforms, make sure you uncomment the CALICO_IPV4POOL_CIDR variable in the manifest and set it to the same value as your chosen pod CIDR.
3. Customize the manifest as necessary
4. Apply the manifest using the following command

```
$ kubectl apply -f calico.yaml
```

Paste it in the master machine:

```
nodes      NotReady   <none>           5s    v1.25.0
[root@master ~]# curl https://projectcalico.docs.tigera.io/manifests/calico.yaml -O
% Total % Received % Xferd Average Speed Time Time Current
          Dload Upload Total Spent Left Speed
100 217k 100 217k 0 0 954k 0 --:--:-- --:--:-- 956k
[root@master ~]#
```

View the calico.yaml file

```
calico.yaml
```

Now apply with this kubectl:

```
[root@master ~]# kubectl apply -f calico.yaml
configmap/calico-config created
customresourcedefinition.apixensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/irpervations.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apixensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
serviceaccount/calico-node created
deployment.apps/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
Warning: policy/vibeta1 PodDisruptionBudget is deprecated in v1.21+, unavailable in v1.25+; use policy/v1 PodDisruptionBudget
poddisruptionbudget.policy/calico-kube-controllers created
```

What is –w? Kubernetes runs your workload by placing containers into Pods to run on Nodes. A node may be a virtual or physical machine, depending on the cluster. Each node is managed by the control plane and contains the services necessary to run Pods.

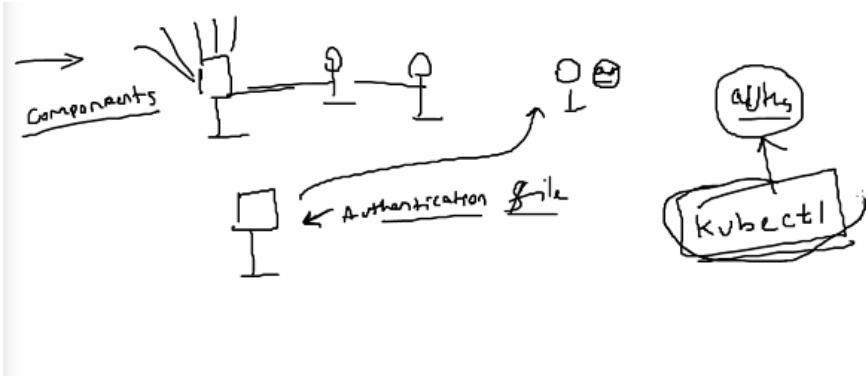
```
[root@master ~]# kubectl get nodes -w
NAME     STATUS    ROLES          AGE      VERSION
master   NotReady control-plane,master 5m51s   v1.23.6
node1   NotReady <none>           2m3s    v1.23.6
node2   NotReady <none>           2m7s    v1.23.6
node3   NotReady <none>           118s   v1.23.6
node4   NotReady <none>           115s   v1.23.6
node5   NotReady <none>           99s    v1.23.6
node3   NotReady <none>           2m2s    v1.23.6
node2   Ready     <none>           2m12s   v1.23.6
node2   Ready     <none>           2m12s   v1.23.6
```

Now everything is in ready state:

```
[root@master ~]# kubectl get nodes
NAME     STATUS    ROLES          AGE      VERSION
master   Ready     control-plane,master 6m9s    v1.23.6
node1   Ready     <none>           2m21s   v1.23.6
node2   Ready     <none>           2m25s   v1.23.6
node3   Ready     <none>           2m16s   v1.23.6
node4   Ready     <none>           2m13s   v1.23.6
node5   Ready     <none>           117s    v1.23.6
[root@master ~]#
```

Created the cluster but everyone can't access it. Only the creator can access it for now.

How to access the cluster:



Using kubectl command:

```
cat <<EOF >/etc/yum.repos.d/kube.repo
[kube]
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
gpgcheck=0
EOF
```

Link: <https://github.com/mdhack0316/k8s>

Screenshot of a GitHub repository page for 'mdhack0316/k8s'. The page contains instructions for setting up a Kubernetes cluster. It includes sections for disabling swap, installing Docker and kubeadm, handling missing kubeadm in repos, configuring yum, and changing docker cgroup driver.

```
cat <<EOF >/etc/yum.repos.d/kube.repo
[kube]
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
gpgcheck=0
EOF
```

You may have to change docker cgroup driver from overlay to Systemd
(since kubernetes 1.20)

```
cat <<K >/etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
```

Install kubectl using this command: yum install kubectl -y

```
[root@master ~]# yum install kubelet -y
```

Read the authentication file from home locations

```
[root@master ~]# pwd
/root
[root@master ~]# cd .kube/
[root@master .kube]# s
bash: s: command not found
[root@master .kube]# ls
cache config
[root@master .kube]# cat con
[root@master .kube]# cat config
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: LS0tLS1CRUdJTiBDRVJUSUZJQ0FUR
kFRc0ZBREFWTViJnd0VRWURWUVFERXdwcmRXSmwKY201bGRHVnpNQjRYRFRJel
WzEzLmQzLmMzLmYzLmVzLmZzLmPzLmTzLmHzzLmKzzLmUzzLmXzzLmYzzLm
kFRc0ZBREFWTViJnd0VRWURWUVFERXdwcmRXSmwKY201bGRHVnpNQjRYRFRJel
WzEzLmQzLmMzLmYzLmVzLmZzLmPzLmTzLmHzzLmKzzLmUzzLmXzzLmYzzLm
```

Mean time creator(trainer) is trying to provide us the authentication file

```
[root@master .kube]# cp config    /var/www/html/
[root@master .kube]# cd
[root@master ~]# chmod -R 777 /var/www/html/
[root@master ~]# cd /var/www/html/
[root@master html]# ls
config
[root@master html]# vim config
```

Help to get you trainer credentials in our machine:

```
yum install wget -y ; wget 3.137.37.209/config ; mkdir $HOME/.kube/ ; cp config $HOME/.kube/ ; kubectl get nodes
```

Verify using this command: `kubectl get nodes`

```
[root@master html]# kubectl get nodes
NAME      STATUS    ROLES          AGE     VERSION
master    Ready     control-plane,master   113m   v1.23.6
node1     Ready     <none>           109m   v1.23.6
node2     Ready     <none>           109m   v1.23.6
node3     Ready     <none>           109m   v1.23.6
node4     Ready     <none>           109m   v1.23.6
node5     Ready     <none>           109m   v1.23.6
[root@master html]#
```

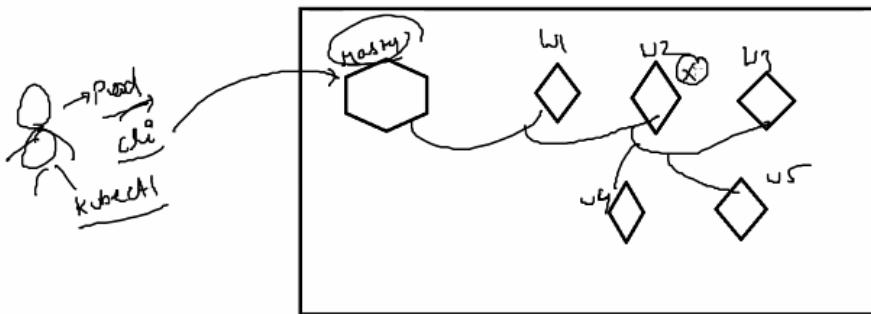
Start the httpd service:

```
[root@master html]# systemctl start httpd  
[root@master html]#
```

What is the need of bash completion? Automation

```
kubectl completion bash > /etc/bash_completion.d/kubectl
```

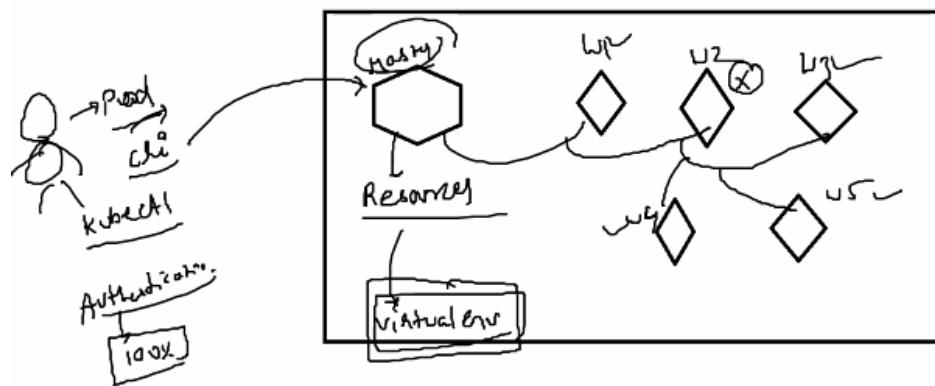
bash



If the worker 2 goes down, where the process or container will go? Anywhere!

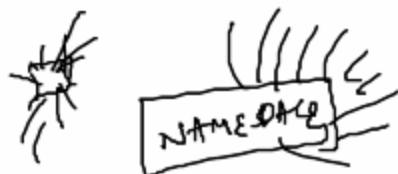
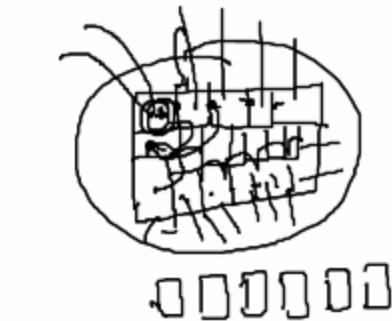
In Kubernetes – we have one great resource - what is that?

Virtual environment – whenever we create a cluster – it will be visible to everyone.



Example: Take my home – 25 rooms – you have a lock to lock the rooms. I have participants and assign to each room. Every participant is in each room.

25



~~26~~

If I don't lock, each participant can see what they are doing in their room.

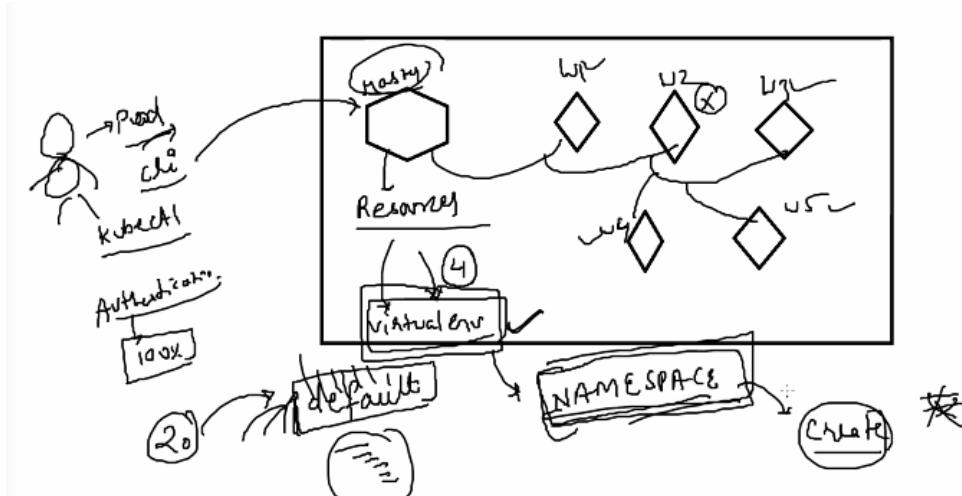
If I lock, then each participant can't see what other participant are doing?

Now, I have given access to everyone where everyone can see what others are doing.

Now I will try to create a personal room – You can create something.

That virtual environment in Kubernetes is called NAMESPACE

If there are multiple developers are there, then there won't be any mix match in the environment. Everyone will have flexibility to develop their own code in their space.



What NAMESPACE is ? It's a resources

Can we verify how many resource we have – Total 4

```
[root@master ~]# kubectl get namespace
NAME        STATUS   AGE
default     Active   124m
kube-node-lease Active  125m
kube-public  Active   125m
kube-system  Active   125m
[root@master ~]#
```

Why default one is taken? The authentication we shared – there is a context name is kuberenetes – this by default using “deafult” as NAMESPACE.

```
kube-system    ACTIVE  125m
[root@master ~]# kubectl get pod
No resources found in default namespace.
[root@master ~]#
```

But now I want to change the NAMESPACE:

```
No resources found in default namespace
[root@master ~]# kubectl create namespace mayank
namespace/mayank created
[root@master ~]#
```

Now you can see mayank namespace is created:

```
[root@master ~]# kubectl get ns
NAME        STATUS   AGE
default     Active   126m
kube-node-lease Active  126m
kube-public  Active   126m
kube-system  Active   126m
mayank      Active   6s
[root@master ~]#
```

Get pods:

```
mayank    ACTIVE  6s
[root@master ~]# kubectl get pods
No resources found in default namespace.
[root@master ~]#
```

You haven't told to the context to use mayank as a namespace instead of default

How to change the context?

```
[root@master ~]# kubectl config
current-context delete-user      get-users      set-cluster      unset
delete-cluster  get-clusters     rename-context  set-context     use-context
delete-context   get-contexts    set           set-credentials  view
[root@master ~]# kubectl config
```

Get context

```
[root@master ~]# kubectl config get-contexts
CURRENT  NAME          CLUSTER      AUTHINFO      NAMESPACE
*        kubernetes-admin@kubernetes  kubernetes  kubernetes-admin
[root@master ~]#
```

Set context

```
*      kubernetes-admin@kubernetes  kubernetes  kubernetes-admin
[root@master ~]# kubectl config set-context kubernetes-admin@kubernetes --namespace mayank
Context "kubernetes-admin@kubernetes" modified.
[root@master ~]#
```

Get context

```
[root@master ~]# kubectl config get-contexts
CURRENT  NAME          CLUSTER      AUTHINFO      NAMESPACE
*        kubernetes-admin@kubernetes  kubernetes   kubernetes-admin  mayank
[root@master ~]# kubectl get pods
No resources found in mayank namespace.
[root@master ~]#
```

Get pods to verify it:

```
No resources found in mayank namespace.
[root@master ~]# kubectl get ns
NAME      STATUS  AGE
default   Active  128m
kube-node-lease  Active  128m
kube-public   Active  128m
kube-system    Active  128m
mayank       Active  2m10s
[root@master ~]# kubectl get pods
No resources found in mayank namespace.
[root@master ~]#
```

Delete the namespace:

```
[root@master ~]# kubectl delete ns manish
namespace "manish" deleted
[
```

Commands shared by trainer:

- 65 kubectl completion bash > /etc/bash_completion.d/kubectl
- 66 bash
- 67 history
- 68 kubectl get namespace
- 69 kubectl get pod
- 70 vim .kube/config
- 71 kubectl get pod
- 72 kubectl create namespace mayank
- 73 kubectl get bs
- 74 kubectl get ns
- 75 kubectl get pods
- 76 bash
- 77 history

2 kubectl config get-contexts

3 kubectl config set-context kubernetes-admin@kubernetes --namespace mayank

4 kubectl config get-contexts

5 kubectl get pods

6 kubectl get ns

7 kubectl get pods

Smallest resources in the Kubernetes? POD

What pod will create? Container

How to check pods? Kubectl get pods

```
[root@master ~]# kubectl get pods
No resources found in mayank namespace.
[root@master ~]#
```

How to create the pod? Declarative method

CLI – Command line Interface

YAML – Yet another markup language

```
[root@master ~]# kubectl run test1 --image httpd
pod/test1 created
[root@master ~]# kube
```

```
pod/test1 created
[root@master ~]# kubectl get pods
NAME READY STATUS RESTARTS AGE
test1 0/1 ContainerCreating 0 4s
[root@master ~]#
```

```
[root@master ~]# kubectl get pods
NAME READY STATUS RESTARTS AGE
test1 1/1 Running 0 15s
```

```
[root@master ~]# kubectl get nodes
NAME STATUS ROLES AGE VERSION
master Ready control-plane,master 136m v1.23.6
node1 Ready <none> 132m v1.23.6
node2 Ready <none> 132m v1.23.6
node3 Ready <none> 132m v1.23.6
node4 Ready <none> 132m v1.23.6
node5 Ready <none> 132m v1.23.6
[root@master ~]#
```

Which machine it is running?

```
node5 Ready <none> 132m v1.23.6
[root@master ~]# kubectl get pods -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED-NODE READINESS GATES
test1 1/1 Running 0 47s 192.168.0.65 node4 <none> <none>
[root@master ~]#
```

In machine 4?

```
[root@node4 ~]# docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
58b1049bdff6 httpd "httpd-foreground" 59 seconds ago Up 58 seconds k8s_test1_te
st1_mayank_24a71837-c5a3-4599-899a-bed00c977989_0
6f3e30a01bbf k8s.gcr.io/pause:3.6 "/pause"
About a minute ago Up About a minute k8s_POD_test
1_mayank_24a71837-c5a3-4599-899a-bed00c977989_0
54baaf59307d calico/node "start-unit" 2 hours ago Up 2 hours k8s_calico-n
ode_calico-node-wsmv2_kube-system_9f653e47-986a-49f2-a662-7709d0ef6a89_0
4e0e5bf645e5 k8s.gcr.io/pause:3.6 "/pause"
2 hours ago Up 2 hours k8s_POD_calic
o-node-wsmv2_kube-system_9f653e47-986a-49f2-a662-7709d0ef6a89_0
185631011c21 k8s.gcr.io/kube-proxy "/usr/local/bin/kube_ 2 hours ago Up 2 hours k8s_kube-pro
xy_kube-proxy-8kqdn_kube-system_73cb394e-0608-48a1-9d5f-8d9ef630407c_0
fd0a022aa6f3 k8s.gcr.io/pause:3.6 "/pause"
2 hours ago Up 2 hours k8s_POD_kube
[root@node4 ~]#
```

Intentionally delete it

```
[root@node4 ~]# docker rm -f k8s_test1_te st1_mayank_24a71837-c5a3-4599-899a-bed00c977989_0
k8s_test1_te st1_mayank_24a71837-c5a3-4599-899a-bed00c977989_0
[root@node4 ~]#
```

Restart:

```
[root@master ~]# kubectl get pods -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
test1 1/1 Running 0 47s 192.168.0.65 node4 <none> <none>
[root@master ~]# kubectl get pods -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
test1 1/1 Running 1 99s 192.168.0.65 node4 <none> <none>
[root@master ~]#
```

We can't do this, as it is not in master node. As it is in different network.

```
test1 ✘ ✘ ✘ Running ✘ 99s 192.168.0.65 node4 <none> <none>
[root@master ~]# curl 192.168.0.65
<html><body><h1>It works!</h1></body></html>
[root@master ~]#
```

To know the information about the pod:

```
[root@master ~]# kubectl describe pod test1
Name: test1
Namespace: mayank
Priority: 0
Node: node4/172.31.5.45
Start Time: Tue, 26 Apr 2022 09:27:13 +0000
Labels: run=nginx
Annotations: cni.projectcalico.org/containerID: 6f3e30a01bbf4c3477bafe0431d1eda65d2638123f55eff536b3f9b5f29
cni.projectcalico.org/podIP: 192.168.0.65/32
cni.projectcalico.org/podIPs: 192.168.0.65/32
Status: Running
IP: 192.168.0.65
IPs:
  IP: 192.168.0.65
Containers:
```

Verification:

```
[root@master ~]# kubectl get pods
NAME READY STATUS RESTARTS AGE
test1 1/1 Running 1 3m34s
[root@master ~]# curl 192.168.0.65
<html><body><h1>It works!</h1></body></html>
[root@master ~]#
```

Delete the pod:

```
[root@master ~]# kubectl delete pod test1
pod "test1" deleted
[root@master ~]#
```

How to create a pod using yaml file?

Create a folder

```
[root@master ~]# ls
[root@master ~]# mkdir code
[root@master ~]# cd code
[root@master code]# ls
[root@master code]#
```

Create and edit the yaml file:

```
[root@master code]# ls
[root@master code]# vim pod.yaml
```

Api-resources:

```
[root@master code]# kubectl api-resources
NAME          SHORTNAMES   APIVERSION      NAMESPACED   KIND
bindings      v1           v1              true        Binding
componentstatuses   cs          v1              false       ComponentStatus
configmaps     cm          v1              true        ConfigMap
endpoints     ep          v1              true        Endpoints
events        ev          v1              true        Event
limits        limits      v1              true        LimitRange
namespaces    ns          v1              false       Namespace
nodes         no          v1              false       Node
persistentvolumeclaims   pvc         v1              true        PersistentVolumeClaim
pvc           v1           v1              false       PersistentVolume
pods          po          v1              true        Pod
podtemplates   pt          v1              true        PodTemplate
replicationcontrollers   rc          v1              true        ReplicationController
secret         quota       v1              true        ResourceQuota
secrets        v1           v1              true        Secret
serviceaccounts   sa          v1              true        ServiceAccount
services        svc         v1              true        Service
mutatingwebhookconfigurations   admissionregistration.k8s.io/v1  false       MutatingWebhookConfiguration
validationwebhookconfigurations   admissionregistration.k8s.io/v1  false       ValidatingWebhookConfiguration
```

Yaml is a decalartive language:

You can kind the apiversion and kind from above

Metadata – some information about the pod

Intention – spacing should be proper – 2 or 3 spaces – same should be followed everywhere

Specification: Write about containers section: name, how many images, We can give multiple values. (hyphen is the separator)

```
apiVersion: v1
kind: Pod
metadata:
  name: FirstPod
spec:
  containers:
    - name: cont1
      image: nginx
```

-- INSERT --

Create the pod using yaml file:

```
apiVersion: v1
kind: Pod
metadata:
  name: FirstPod
spec:
  containers:
    - name: cont1
      image: nginx
```

Problem here is the pod name should not be in caps. Make it lower "firstpod"

```
[root@master code]# kubectl create -f pod.yaml
pod/firstpod created
[root@master code]#
```

Get pod:

```
[root@master code]# kubectl get pods
NAME      READY  STATUS    RESTARTS   AGE
firstpod  1/1    Running   0          7s
[root@master code]#
```

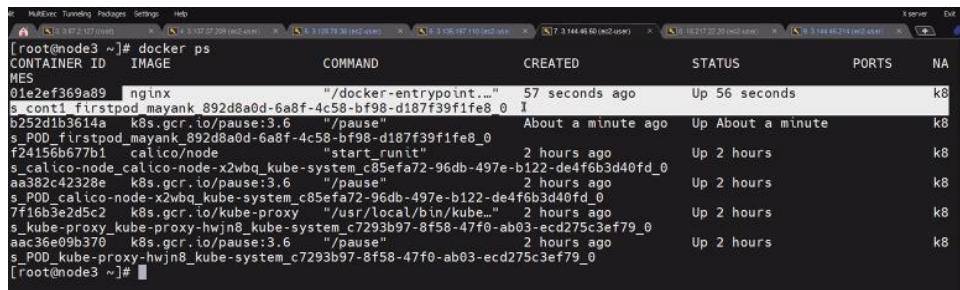
Get pods wide:

```
[root@master code]# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE   READINESS GATES
firstpod  1/1     Running   0          24s    192.168.0.193   node3  <none>        <none>
```

Curl command:

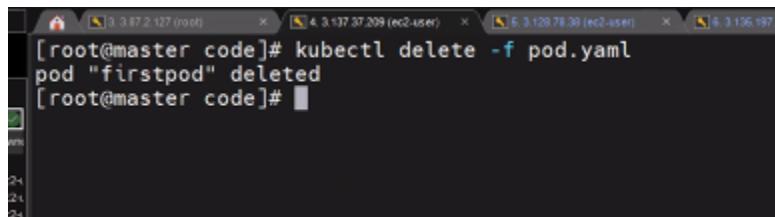
curl 192.168.0.193

You can check by docker ps:



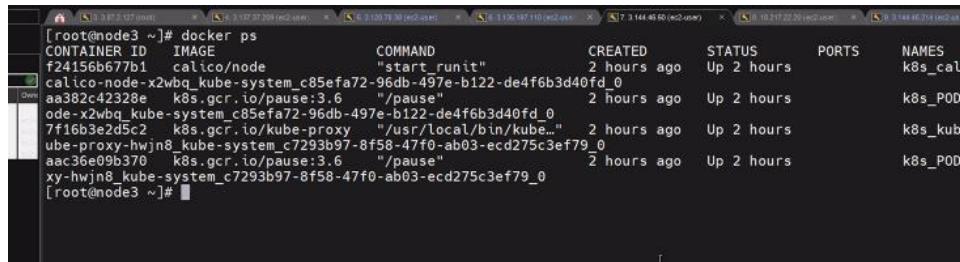
```
[root@node3 ~]# docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS           NAMES
01e2ef369a89        nginx              "/docker-entrypoint..."   57 seconds ago     Up 56 seconds       k8s
b252d1b3614a        k8s.gcr.io/pause:3.6   "pause"           About a minute ago   Up About a minute   k8s_POD_firstpod_mayan...
s_POD_firstpod_mayan..._892d8a0d-6a8f-4c58-bf98-d187f39f1fe8_0
f24156b677b1        calico/node          "start_runit"      2 hours ago       Up 2 hours         k8s
s_calico-node_calico-node-x2wbq_kube-system_c85efa72-96db-497e-b122-de4f6b3d40fd_0
aa382c42328e        k8s.gcr.io/pause:3.6   "pause"           2 hours ago       Up 2 hours         k8s_POD_calico-node-x2...
7f16b3e2d5c2        k8s.gcr.io/kube-proxy   "/usr/local/bin/kube..."  2 hours ago       Up 2 hours         k8s_kube-proxy_kube-proxy-hw...
aac36e09b370        k8s.gcr.io/pause:3.6   "pause"           2 hours ago       Up 2 hours         k8s_POD_kube-proxy_hw...
[root@node3 ~]#
```

Delete the pod:



```
[root@master code]# kubectl delete -f pod.yaml
pod "firstpod" deleted
[root@master code]#
```

In machine 3 – it is not visible:



```
[root@node3 ~]# docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS           NAMES
calico-node-x2wbq_kube-system_c85efa72-96db-497e-b122-de4f6b3d40fd_0
aa382c42328e        k8s.gcr.io/pause:3.6   "pause"           2 hours ago       Up 2 hours         k8s_POD_calico-node-x2...
7f16b3e2d5c2        k8s.gcr.io/kube-proxy   "/usr/local/bin/kube..."  2 hours ago       Up 2 hours         k8s_kube-proxy_kube-proxy-hw...
aac36e09b370        k8s.gcr.io/pause:3.6   "pause"           2 hours ago       Up 2 hours         k8s_POD_kube-proxy_hw...
xy-hwjn8_kube-system_c7293b97-8f58-47f0-ab03-e...
[root@node3 ~]#
```

5 kubectl get pods

6 kubectl get ns

7 kubectl get pods

8 history

9 kubectl get ns

10 kubectl delete ns manish

11 kubectl get ns

12 kubectl get pods

```
13 kubectl get namespaces
14 kubectl get pods
15 kubectl run test1 --image httpd
16 kubectl get pods
17 kubectl get nodes
18 kubectl get pods -o wide
19 curl 192.168.0.65
20 kubectl describe pod test1
21 kubectl get pods
22 curl 192.168.0.65
23 kubectl delete pod test1
24 kubectl get pods
25 kubectl get pods --all-namespaces
26 kubectl get pods
27 mkdir code
28 cd code
29 ls
30 vim pod.yaml
31 kubectl api-resources
32 vim pod.yaml
33 kubectl get pods
34 kubectl create -f pod.yaml
35 kubectl create -f pod.yaml
36 vim pod.yaml
37 kubectl create -f pod.yaml
38 kubectl get pods
39 kubectl get pods -o wide
40 curl 192.168.0.193
41 kubectl describe pod firstpod
42 kubectl delete -f pod.yaml
```

43 history

Yaml file:

```
apiVersion: v1
kind: Pod
metadata:
  name: firstpod
spec:
  containers:
    - name: cont1
      image: nginx
```

TASK:

<http://54.159.140.24/kube/task1.php>

Q.1 Create A namespace with your own name.
Q.2 Modify your context to use your name namespace.
Q.3 Create a Pod using httpd image and name of POD should be pod1.
Q.4 Create a Pod using nginx image and name of POD should be pod2.
Q.5 Create a Pod using yaml method , and use nginx image and name of POD should be yamlpod1.
Q.6 Create a Pod using yaml method , and use mdhack/myserver image and name of POD should be yamlpod2.

Task:

Question:

Q.1 Create a namespace with your own name.

Get namespace:

```
[root@vishali ~]# kubectl get namespace
NAME          STATUS   AGE
default       Active   153m
kube-node-lease Active  154m
kube-public   Active   154m
kube-system   Active   154m
manish        Active   3m4s
manoj         Active   6s
mayank        Active   27m
praval        Active   10s
sampat        Active   33s
sankalp       Active   36s
sashi         Active   51s
vaishnavi    Active   48s
```

Get pod:

```
[root@vishali ~]# kubectl get pod
No resources found in default namespace.
```

Create the namespace:

```
[root@vishali ~]# kubectl create namespace vishali
namespace/vishali created
```

Verify it:

```
[root@vishali ~]# kubectl get ns
NAME          STATUS   AGE
alpana        Active   3s
default       Active   154m
kaustubh     Active   28s
kube-node-lease Active   154m
kube-public   Active   154m
kube-system   Active   154m
manish        Active   4m
manoj         Active   62s
mayank        Active   28m
praval        Active   66s
sampat        Active   89s
sankalp       Active   92s
sashi         Active   107s
shweta        Active   40s
vaishnavi    Active   104s
vishali      Active   9s
```

But pods has got default namespace, we need to change the context

```
[root@vishali ~]# kubectl get pods
No resources found in default namespace.
```

Q.2 Modify your context to use your name namespace.

Get context:

```
[root@vishali ~]# kubectl config get-contexts
CURRENT  NAME          CLUSTER      AUTHINFO      NAMESPAC
*        kubernetes-admin@kubernetes  kubernetes  kubernetes-admin
```

Set the context:

```
[root@vishali ~]# kubectl config set-context kubernetes-admin@kubernetes --name=vishali
Context "kubernetes-admin@kubernetes" modified.
```

Again, get the context to verify:

```
context "kubernetes-admin@kubernetes" modified.
[root@vishali ~]# kubectl config get-contexts
CURRENT  NAME          CLUSTER      AUTHINFO      NAMESPAC
*        kubernetes-admin@kubernetes  kubernetes  kubernetes-admin  vishali
```

Get the pods:

```
[root@vishali ~]# kubectl get pods
No resources found in vishali namespace.
[root@vishali ~]# kubectl get ns
NAME          STATUS   AGE
alpana        Active   2m31s
default       Active   157m
kaustubh     Active   2m56s
kube-node-lease Active   157m
kube-public   Active   157m
kube-system   Active   157m
manish        Active   6m28s
manoj         Active   3m30s
mayank        Active   30m
praval        Active   3m34s
sampat        Active   3m57s
sangamesh    Active   2m22s
sankalp       Active   4m
sashi         Active   4m15s
shweta        Active   3m8s
vaishnavi    Active   4m12s
vishali       Active   2m37s
```

Q.3 Create a Pod using httpd image and name of POD should be pod1.

```
[root@vishali ~]# kubectl run pod1 --image httpd
pod/pod1 created
[root@vishali ~]# kubectl get pods
NAME    READY  STATUS   RESTARTS   AGE
pod1   1/1    Running   0          7s
[root@vishali ~]# kubectl get nodes
NAME      STATUS   ROLES      AGE   VERSION
master   Ready    control-plane,master 164m  v1.23.6
node1    Ready    <none>    160m  v1.23.6
node2    Ready    <none>    160m  v1.23.6
node3    Ready    <none>    160m  v1.23.6
node4    Ready    <none>    160m  v1.23.6
node5    Ready    <none>    159m  v1.23.6
[root@vishali ~]# kubectl get pods -o wide
NAME    READY  STATUS   RESTARTS   AGE   IP           NODE   NOMINATED NODE
      READINESS GATES
pod1   1/1    Running   0          39s   192.168.0.133  node1  <none>
```

Q.4 Create a Pod using nginx image and name of POD should be pod2.

```
[root@vishali ~]# kubectl run pod2 --image nginx
pod/pod2 created
[root@vishali ~]# kubectl get pods
NAME    READY  STATUS   RESTARTS  AGE
pod1   1/1    Running  0          4m41s
pod2   1/1    Running  0          10s
[root@vishali ~]# kubectl get pods -o wide
NAME    READY  STATUS   RESTARTS  AGE      IP           NODE  NOMINATED NO
DE    READINESS GATES
pod1   1/1    Running  0          4m58s   192.168.0.133  node1 <none>
<none>
pod2   1/1    Running  0          27s     192.168.0.201  node3 <none>
<none>
[root@vishali ~]#
```

Q.5 Create a Pod using yaml method and use nginx image and name of POD should be yamlpod1.

Create a folder:

```
[root@vishali ~]# mkdir code
[root@vishali ~]# cd code
[root@vishali code]# vim pod.yaml
```

Yaml file:

```
[root@vishali code]# cat pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: yamlpod1
spec:
  containers:
    - name: cont1
      image: nginx
```

Create the pod using that yaml file:

```
[root@vishali code]# kubectl create -f pod.yaml
pod/yamlpod1 created
```

Q.6 Create a Pod using yaml method and use mdhack/myserver image and name of POD should be yamlpod2.

Yaml file:

```
[root@vishali code]# cat pod2.yaml
apiVersion: v1
kind: Pod
metadata:
  name: yamlpod2
spec:
  containers:
    - name: cont2
      image: mdhack/myserver
```

Create the pod using that yaml file:

```
[root@vishali code]# kubectl create -f pod2.yaml
pod/yamlpod2 created
```

Kubectl get pods:

```
[root@vishali code]# kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
pod1      1/1     Running   0          15m
pod2      1/1     Running   0          10m
yamlpod1  1/1     Running   0          5m34s
yamlpod2  1/1     Running   0          83s
```

Kubectl get pods -o wide:

```
yamlpod2  1/1     Running   0          83s
[root@vishali code]# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATE
D NODE   READINESS GATES
pod1      1/1     Running   0          15m   192.168.0.133 node1 <none>
      <none>
pod2      1/1     Running   0          11m   192.168.0.201 node3 <none>
      <none>
yamlpod1  1/1     Running   0          5m59s  192.168.0.203 node3 <none>
      <none>
yamlpod2  1/1     Running   0          108s  192.168.0.140 node1 <none>
      <none>
[root@vishali code]#
```

If there is no pod entity, what is the other methods? Virtual environment.



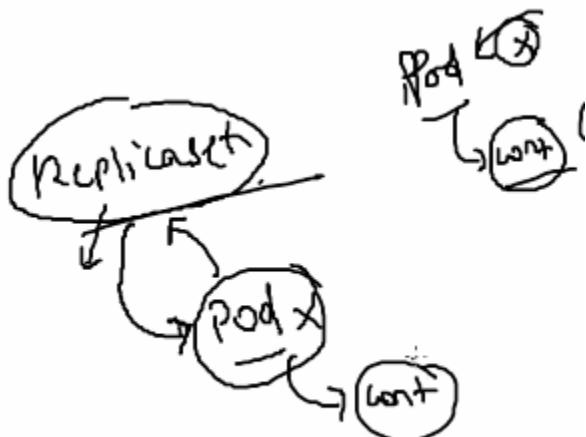
If we create a docker container and delete it? Will it come back? No

Same thing is with pod. It will never come back.

But as the part of high availability part, it should come back. We need to create another feature – resource called deployment. Without down time it can be changed

Pod is an entity, but we are not going to be created automatically.

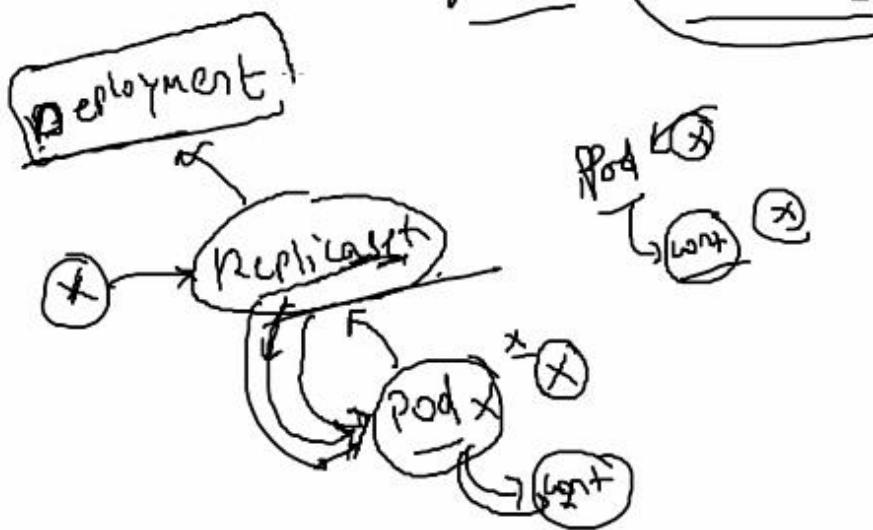
We have replica set – help us to create pods automatically:



Pod will create container.

If you delete the pod, container will delete. What replica will ensure? It will ensure pod will automatically come back.

Replica we can use it but it can't have one feature? What is that?



When deployment is created, then replica set will be created. If the replica is deleted, then deployment is created automatically. You can create n number of replica set. One replica set can have multiple pods.

Deployment: It can be created by cli and yaml.

```
[root@master code]# kubectl get all -n default
NAME           READY   STATUS    RESTARTS   AGE
pod/firstpod   1/1     Running   0          44m
pod/pod1       1/1     Running   0          53m
pod/pod2       0/1     ImagePullBackOff 0          52m
pod/pod22      1/1     Running   0          43m
pod/pod3       0/1     ImagePullBackOff 0          51m
pod/yamlpod1   1/1     Running   0          42m
pod/yamlpod2   1/1     Running   0          41m

NAME            TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/kubernetes   ClusterIP   10.96.0.1   <none>        443/TCP   3h39m
[root@master code]# kubecget ns
bash: kubecget: command not found
```

To know the deployments created:

```
[root@master code]# kubectl get deployments.apps
No resources found in mayank namespace.
[root@master code]#
```

get deployments: (not necessary to use .apps)

```
[root@master code]# kubectl get deployments.apps
No resources found in mayank namespace.
[root@master code]# kubectl get deployments
No resources found in mayank namespace.
[root@master code]#
```

get deployments, pods, replicas

```
[root@master code]# kubectl get deployments,pods,rs
No resources found in mayank namespace.
[root@master code]#
```

get all:

```
No resources found in mayank namespace.
[root@master code]# kubectl get all
No resources found in mayank namespace.
[root@master code]#
```

Using CLI:

create deployment:

```
[root@master code]# kubectl create deployment mayank --image=nginx
deployment.apps/mayank created
[root@master code]#
```

get all:

get all -o wide:

```
[root@master code]# kubectl get all -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATE
ES
pod/mayank-9c4f4b955-ctdcl  1/1    Running   0          22s    192.168.0.31  node5  <none>        <none>
NAME          READY   UP-TO-DATE   AVAILABLE   AGE     CONTAINERS   IMAGES   SELECTOR
deployment.apps/mayank      1/1    1          1          22s    nginx        nginx    app=mayank
NAME          DESIRED  CURRENT   READY     AGE     CONTAINERS   IMAGES   SELECTOR
replicaset.apps/mayank-9c4f4b955  1       1       1       22s    nginx        nginx    app=mayank,pod-template-ha
sh=9c4f4b955
[root@master code]#
```

Verify:

```
[root@master code]# curl 192.168.0.31
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
```

Delete the pod – it will create the pod automatically

```
[root@master code]# kubectl delete pod mayank-9c4f4b955-ctdcl --force
warning: Immediate deletion does not wait for confirmation that the running resource has been terminated.
It may continue to run on the cluster indefinitely.
pod "mayank-9c4f4b955-ctdcl" force deleted
[root@master code]# kubectl get all
NAME           READY   STATUS      RESTARTS   AGE
pod/mayank-9c4f4b955-lcb48  0/1   ContainerCreating   0          2s
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mayank  0/1     1           0          89s
NAME           DESIRED  CURRENT    READY      AGE
replicaset.apps/mayank-9c4f4b955  1       1           0          89s
[root@master code]#
```

verify it:

```
[root@master code]# kubectl get all -o wide
NAME           READY   STATUS      RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
pod/mayank-9c4f4b955-lcb48  1/1   Running   0          18s  192.168.0.208  node3  <none>        <none>
NAME           READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES   SELECTOR
deployment.apps/mayank  1/1     1           1          105s  nginx        nginx   app=mayank
NAME           DESIRED  CURRENT    READY      AGE   CONTAINERS   IMAGES   SELECTOR
replicaset.apps/mayank-9c4f4b955  1       1           1          105s  nginx        nginx   app=mayank,pod-template-h
ash=9c4f4b955
[root@master code]# kubectl delete rs mayank-9c4f4b955
```

Delete replica – automatically created:

```
[root@master code]# kubectl delete rs mayank-9c4f4b955
replicaset.apps "mayank-9c4f4b955" deleted
[root@master code]# kubectl get all -o wide
NAME           READY   STATUS      RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
pod/mayank-9c4f4b955-d5k6x  0/1   ContainerCreating   0   2s  <none>        node5  <none>        <none>
NAME           READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES   SELECTOR
deployment.apps/mayank  0/1     1           0          2m12s  nginx        nginx   app=mayank
NAME           DESIRED  CURRENT    READY      AGE   CONTAINERS   IMAGES   SELECTOR
replicaset.apps/mayank-9c4f4b955  1       1           0          2s   nginx        nginx   app=mayank,pod-template-h
ash=9c4f4b955
[root@master code]#
```

Verify:

```
[root@master code]# curl 192.168.0.32
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
[root@master code]#
```

get deployment:

```
[root@master code]# kubectl get deployments.apps
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mayank    1/1     1           1           3m36s
[root@master code]# kubectl edit ■
```

Edit the deployment:

```
[root@master code]# kubectl edit   deployments.apps  mayank ■
```

```
mayank   1/1   1           1           3m36s
[root@master code]# kubectl edit   deployments.apps  mayank
Edit cancelled, no changes made.
[root@master code]#
```

If you want to create a deployment - want to get yaml automatically:

Dry run: Test

```
[root@master code]# kubectl get deployments.apps
[root@master code]# kubectl get deployments.apps
NAME    READY  UP-TO-DATE AVAILABLE AGE
mayank  1/1    1           1        3m36s
[root@master code]# kubectl edit deployments.apps mayank
Edit cancelled, no changes made.
(reverse-i-search)`dep': kubectl edit ^Cployments.apps mayank
[root@master code]# kubectl create deployment mayank --image=nginx --dry-run
[root@master code]# kubectl create deployment mayank --image=nginx --dry-run
[0426 10:57:32.386642    3726 helpers.go:598] --dry-run is deprecated and can be replaced with
deployment.apps/mayank created (dry run)
[root@master code]#
```

```
[root@master code]# kubectl create deployment mayank --image=nginx --dry-run -o yaml
[0426 10:57:55.199045    4094 helpers.go:598] --dry-run is deprecated and can be replaced with --dry-run=client.
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mayank
    name: mayank
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mayank
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mayank
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}    }
status: {}
[root@master code]#
```

Create the deployment with dry run:

```
[root@master code]# kubectl create deployment mayank --image=nginx --dry-run -o yaml > deployment.yaml
[0426 10:58:13.634753    4405 helpers.go:598] --dry-run is deprecated and can be replaced with --dry-run=client.
[root@master code]# vim de
```

yaml is completely ready:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mayank
  name: mayank
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mayank
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mayank
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
status: {}
~
~
```

Yaml file:

How to deploy the yaml file:

create command:

```
[root@master code]# vim deployment.yaml
[root@master code]# kubectl create -f deployment.yaml
Error from server (AlreadyExists): error when creating "deployment.yaml": deployments.apps "mayank" already exists
[root@master code]#
```

```
[root@master code]# kubectl create -f deployment.yaml
```

Get all:

```
[root@master code]# kubectl get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/mayank-9c4f4b955-jg5x4          1/1     Running   0          5s
NAME                            READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mayank             1/1     1           1           5s
NAME                            DESIRED   CURRENT   READY   AGE
replicaset.apps/mayank-9c4f4b955  1         1         1         5s
[root@master code]#
```

Delete the pod:

```
replicaset.apps/mayank-9c4f4b955  1         1         1         5s
[root@master code]# kubectl delete pod mayank-9c4f4b955-jg5x4 -fo
error: the path "o" does not exist
[root@master code]# kubectl delete pod mayank-9c4f4b955-jg5x4 --force
warning: Immediate deletion does not wait for confirmation that the running resource has
e may continue to run on the cluster indefinitely.
pod "mayank-9c4f4b955-jg5x4" force deleted
[root@master code]# kubectl delete pod mayank-9c4f4b955-jg5x4 -fo
error: the path "o" does not exist
[root@master code]# kubectl get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/mayank-9c4f4b955-lnsxs          1/1     Running   0          3s
NAME                            READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mayank             1/1     1           1           37s
NAME                            DESIRED   CURRENT   READY   AGE
replicaset.apps/mayank-9c4f4b955  1         1         1         37s
[root@master code]#
```

Get wide:

```
[root@master code]# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
mayank-9c4f4b955-lnsxs              1/1     Running   0          40s  192.168.0.33  node5  <none>        <none>
[root@master code]# curl
```

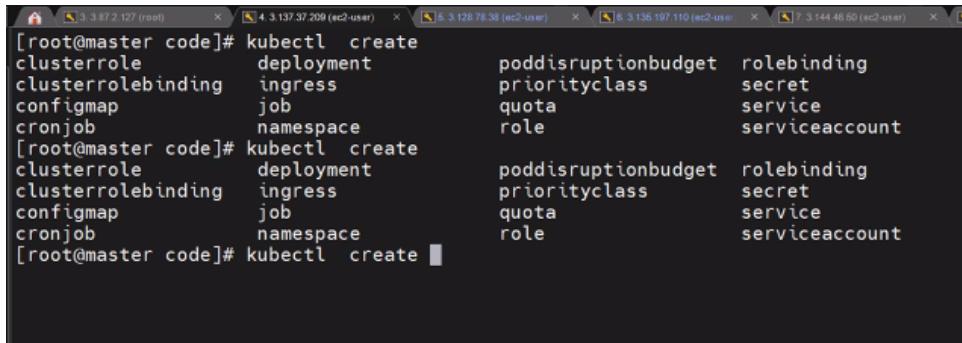
Curl – verify:

```
[root@master code]# curl 192.168.0.33
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
[root@master code]#
```

These many things you can create:



```
[root@master code]# kubectl create deployment poddisruptionbudget rolebinding
clusterrole deployment priorityclass secret
clusterrolebinding ingress priorityclass service
configmap job quota service
cronjob namespace role serviceaccount
[root@master code]# kubectl create deployment poddisruptionbudget rolebinding
clusterrole deployment priorityclass secret
clusterrolebinding ingress priorityclass service
configmap job quota service
cronjob namespace role serviceaccount
[root@master code]# kubectl create [REDACTED]
```

Deployment Commands shared by Trainer:

- 47 kubectl get ns
- 48 kubectl get pods
- 49 kubectl get pods -n default
- 50 kubectl get all -n default
- 51 kubecget ns
- 52 kubectl get ns
- 53 kubectl get ns
- 54 kubectl get deployments.apps
- 55 kubectl get deployments
- 56 kubectl get deployments,pods,rs
- 57 kubectl get all
- 58 kubectl create deployment mayank --image=nginx
- 59 kubectl get all
- 60 kubectl get all -o wide
- 61 curl 192.168.0.31
- 62 kubectl get all
- 63 kubectl delete pod mayank-9c4f4b955-ctdcl --force
- 64 kubectl get all
- 65 kubectl get all -o wide
- 66 kubectl delete rs mayank-9c4f4b955
- 67 kubectl get all -o wide
- 68 curl 192.168.0.32

```
69 kubectl get deployments.apps
70 kubectl edit deployments.apps mayank
71 kubectl create deployment mayank --image=nginx --dry-run
72 kubectl create deployment mayank --image=nginx --dry-run -o yaml
73 kubectl create deployment mayank --image=nginx --dry-run -o yaml > deployment.yaml
74 vim deployment.yaml
75 kubectl create -f deployment.yaml
76 kubectl delete deployments.apps mayank
77 kubectl create -f deployment.yaml
78 kubectl get all
79 kubectl delete pod mayank-9c4f4b955-jg5x4 -fo
80 kubectl delete pod mayank-9c4f4b955-jg5x4 --force
81 kubectl delete pod mayank-9c4f4b955-jg5x4 -fo
82 kubectl get all
83 kubectl get pods -o wide
84 curl 192.168.0.33
85 kubectl describe deployments.apps mayank
86 history
```

Yaml file:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mayank
  name: mayank
spec:
  replicas: 1
  selector:
```

```
matchLabels:  
  app: mayank  
strategy: {}  
template:  
  metadata:  
    creationTimestamp: null  
  labels:  
    app: mayank  
spec:  
  containers:  
    - image: nginx  
      name: nginx  
      resources: {}  
status: {}
```

Task:

<http://54.159.140.24/kube/task2.php>

Q.1 Create A deployment using CLi Method using mdhack/myapache image and deployment name should be test1.

Create the deployment:

```
[root@vishali ~]# kubectl create deployment test1 --image=mdhack/myapache  
deployment.apps/test1 created
```

Verify it:

```
[root@vishali ~]# kubectl get deployments  
NAME      READY   UP-TO-DATE   AVAILABLE   AGE  
test1     1/1     1           1           20s  
[root@vishali ~]#
```

Q.2 Create A deployment using generating yaml Method using mdhack/myserver image and deployment name should be test2

```

[root@vishali ~]# create deployment test2 --image=mdhack/myserver --dry-run -o yaml > deployment.yaml
-bash: create: command not found
[root@vishali ~]# ^C
[root@vishali ~]# kubectl create deployment test2 --image=mdhack/myserver --dry-run -o yaml > deployment.yaml
W0426 11:17:47.731511 10406 helpers.go:598] --dry-run is deprecated and can be replaced with --dry-run=client.
[root@vishali ~]# kubectl create -f deployment.yaml
deployment.apps/test2 created
[root@vishali ~]# kubectl get all
NAME          READY   STATUS    RESTARTS   AGE
pod/pod1      1/1     Running   0          83m
pod/pod2      1/1     Running   0          79m
pod/test1-6db788bdd4-mccwb 1/1     Running   0          4m55s
pod/test2-744cdc6c89-dvpgb 1/1     Running   0          11s
pod/yamlpod1  1/1     Running   0          73m
pod/yamlpod2  1/1     Running   0          69m

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/test1  1/1         1           1          4m55s
deployment.apps/test2  1/1         1           1          11s

NAME          DESIRED   CURRENT   READY   AGE
replicaset.apps/test1-6db788bdd4  1         1         1        4m55s
replicaset.apps/test2-744cdc6c89  1         1         1        11s
[root@vishali ~]#

```

```

[root@vishali ~]# kubectl get pods -o wide
NAME          NOMINATED NODE   READINESS   STATUS    RESTARTS   AGE   IP          NO
DE          Nominated Node   Readiness   Status   Restarts   Age
pod1        <none>           <none>      Running  0          84m  192.168.0.133 no
de1         <none>           <none>      Running  0          80m  192.168.0.201 no
pod2        <none>           <none>      Running  0          80m  192.168.0.150 no
de3         <none>           <none>      Running  0          5m51s 192.168.0.88 no
test1-6db788bdd4-mccwb 1/1     Running   0          5m51s    192.168.0.203 no
de1         <none>           <none>      Running  0          67s   192.168.0.140 no
test2-744cdc6c89-dvpgb 1/1     Running   0          67s   192.168.0.203 no
de4         <none>           <none>      Running  0          74m   192.168.0.203 no
yamlpod1   1/1               Running   0          74m   192.168.0.203 no
de3         <none>           <none>      Running  0          70m   192.168.0.203 no
yamlpod2   1/1               Running   0          70m   192.168.0.203 no
de1         <none>           <none>      Running  0          70m   192.168.0.203 no
[root@vishali ~]#

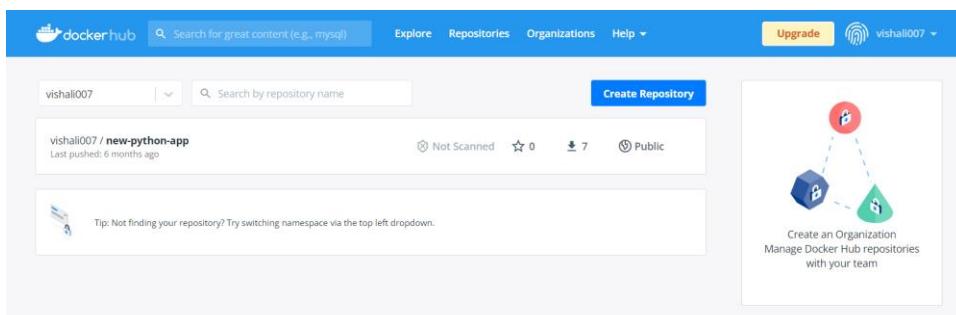
```

Q.3 Create a dockerfile using oraclelinux:8.3 image configure Apache into it with 3 html files into it. and then upload the newly created image on yourdocker hub account. Use that same image while creating the deployment that u have uploaded.

Q.4 Create a deployment using image quay.io/mayank123modi/simple-webapp After creating this expose it on LoadBalancer by CLI method and check everything.

Post Task:

Q.1 create an account on Dockerhub



Q.2 create your own docker image with having 2 diff html files. And find a way to upload those images on your docker hub account

Edit the dockerfile:

```
[root@vishali /]# cd task3
[root@vishali task3]# ls
Dockerfile
[root@vishali task3]# cat Dockerfile
FROM oraclelinux:8.3
RUN yum install httpd -y
WORKDIR /var/www/html/
RUN echo Hello From Dockerfile > index.html
RUN echo Hello from Training file > training.html
CMD ["httpd", "-DFOREGROUND"]
[root@vishali task3]#
```

Docker build:

```
[root@vishali task3]# vim Dockerfile
[root@vishali task3]# docker build -t vishaliimage .
Sending build context to Docker daemon 2.048kB
Step 1/6 : FROM oraclelinux:8.3
8.3: Pulling from library/oraclelinux
dd34f38d274c: Pull complete
Digest: sha256:af3182ee6c1e56f18fc1fecaf638da57d7c47233862e5c32
Status: Downloaded newer image for oraclelinux:8.3
--> 816d99f0bbe8
Step 2/6 : RUN yum install httpd -y
--> Running in 1eab74182ca4
Oracle Linux 8 BaseOS Latest (x86_64) 42 MB/s | 43 M
```

Docker images:

```
[root@vishali task3]# docker images
REPOSITORY      TAG      IMAGE ID      CREATED      SIZE
vishaliimage    latest   60e49d69ed90  15 seconds ago  393MB
newnginx        latest   20a73b7c11ab  5 hours ago   142MB
httpdimage      latest   0bde37eabb5b  5 hours ago   144MB
nginx           latest   fa5269854a5e  6 days ago    142MB
httpd           latest   c30a46771695  6 days ago    144MB
oraclelinux     8.3     816d99f0bbe8  12 months ago  224MB
```

Docker login: using the command docker login

```
[root@vishali task3]# docker login
Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, head over to https://hub.docker.com to create one.
Username: vishali007
Password:
WARNING! Your password will be stored unencrypted in /root/.docker/config.json. Configure a credential helper to remove this warning. See https://docs.docker.com/engine/reference/commandline/login/#credentials-store
Login Succeeded
```

Perform tag and push the image:

```
# docker image tag vishaliimage:latest vishali007/vishaliimage
```

```
# docker image push vishali007/vishaliimage:1.0
```

```
[root@vishali task3]# docker image tag vishaliimage:latest vishali007/vishaliimage:1.0
[root@vishali task3]# docker image push vishali007/vishaliimage:1.0
The push refers to repository [docker.io/vishali007/vishaliimage]
0b5df366ee68: Pushed
daf56a403ed6: Pushed
ea6f8eb24f4b: Pushed
02a3a073ed48: Mounted from library/oraclelinux
1.0: digest: sha256:ccd6c63949f3aa343685623569aa5d50bdabeea4f836ac7a72eae3f5a8d3
02c5 size: 1156
[root@vishali task3]#
```

Docker images are uploaded in the docker hub:

The screenshot shows the Docker Hub interface. At the top, there's a search bar with 'Search for great content' and a button 'Upgrade'. Below the search bar, there are navigation links: Explore, Repositories, Organizations, Help, and a yellow 'Upgrade' button. A dropdown menu is open over the 'Repositories' link. On the left, there's a sidebar with a user icon and the text 'vishali007'. To its right is another search bar labeled 'Search by repository name'. At the top right is a blue 'Create Repository' button. The main content area displays three repository cards for the user 'vishali007':

- vishali007 / postimage**: Last pushed 2 minutes ago. Status: Not Scanned. Stars: 0. Downloads: 0. Public.
- vishali007 / vishaliimage**: Last pushed 15 minutes ago. Status: Not Scanned. Stars: 0. Downloads: 0. Public.
- vishali007 / new-python-app**: Last pushed 6 months ago. Status: Not Scanned. Stars: 0. Downloads: 7. Public.

Q.3 Use the same image that u has uploaded to create deployment in Kubernetes cluster.

Create the deployment:

```
# docker run -itd --name newcont1 vishali007/vishaliimage:1.0
```

```
[root@vishali task3]# docker run -itd --name newcont1 vishali007/postimage:1.0
83381afa0268c57276c244e039d7fed979b439e06a02474263adab4f5bf21bd8
[root@vishali task3]# kubectl create deployment postimagetask3 --image=vishali007/postimage:1.0
deployment.apps/postimagetask3 created
```

Verify the status of the pod:

```
[root@vishali ~]# kubectl get pod
NAME                               READY   STATUS    RESTARTS   AGE
pod1                                1/1    Running   0          125m
pod2                                1/1    Running   0          121m
postimagetask-6f4db4bcd5-vqj7d7    0/1    ImagePullBackOff 0          17m
postimagetask3-db9b9dd8d-p8kf7     1/1    Running   0          11m
posttask-59b9548d4c-qhxs7         0/1    ImagePullBackOff 0          21m
posttask3-7ff7f9b799-g22p2        0/1    ImagePullBackOff 0          28m
test1-6db788bdd4-mccwb           1/1    Running   0          47m
test2-744cdc6c89-dvpgb           1/1    Running   0          42m
yamlpod1                           1/1    Running   0          115m
yamlpod2                           1/1    Running   0          111m
[root@vishali task3]#
```

Get the IP details using the `-o wide` command:

```
92.168.0.140 node1 <none> <none>
[root@vishali task3]# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   I
P          NODE    NOMINATED NODE   READINESS GATES
pod1      1/1    Running   0          114m  1
92.168.0.133 node1 <none> <none>
pod2      1/1    Running   0          109m  1
92.168.0.201 node3 <none> <none>
postimagetask-6f4db4bcd5-vqj7d7 0/1    ImagePullBackOff 0          6m13s  1
92.168.0.218 node3 <none> <none>
postimagetask3-db9b9dd8d-p8kf7  1/1    Running   0          21s   1
92.168.0.52  node2 <none> <none>
posttask-59b9548d4c-qhxs7       0/1    ImagePullBackOff 0          10m   1
92.168.0.50  node2 <none> <none>
posttask3-7ff7f9b799-g22p2     0/1    ImagePullBackOff 0          16m   1
92.168.0.47  node2 <none> <none>
test1-6db788bdd4-mccwb        1/1    Running   0          35m   1
92.168.0.150 node1 <none> <none>
test2-744cdc6c89-dvpgb        1/1    Running   0          30m   1
92.168.0.88  node4 <none> <none>
yamlpod1                          1/1    Running   0          104m  1
92.168.0.203 node3 <none> <none>
yamlpod2                          1/1    Running   0          100m  1
92.168.0.140 node1 <none> <none>
[root@vishali task3]# ^C
```

Pod IP Address: 192.168.0.52

Alternative of how to try these above commands:

